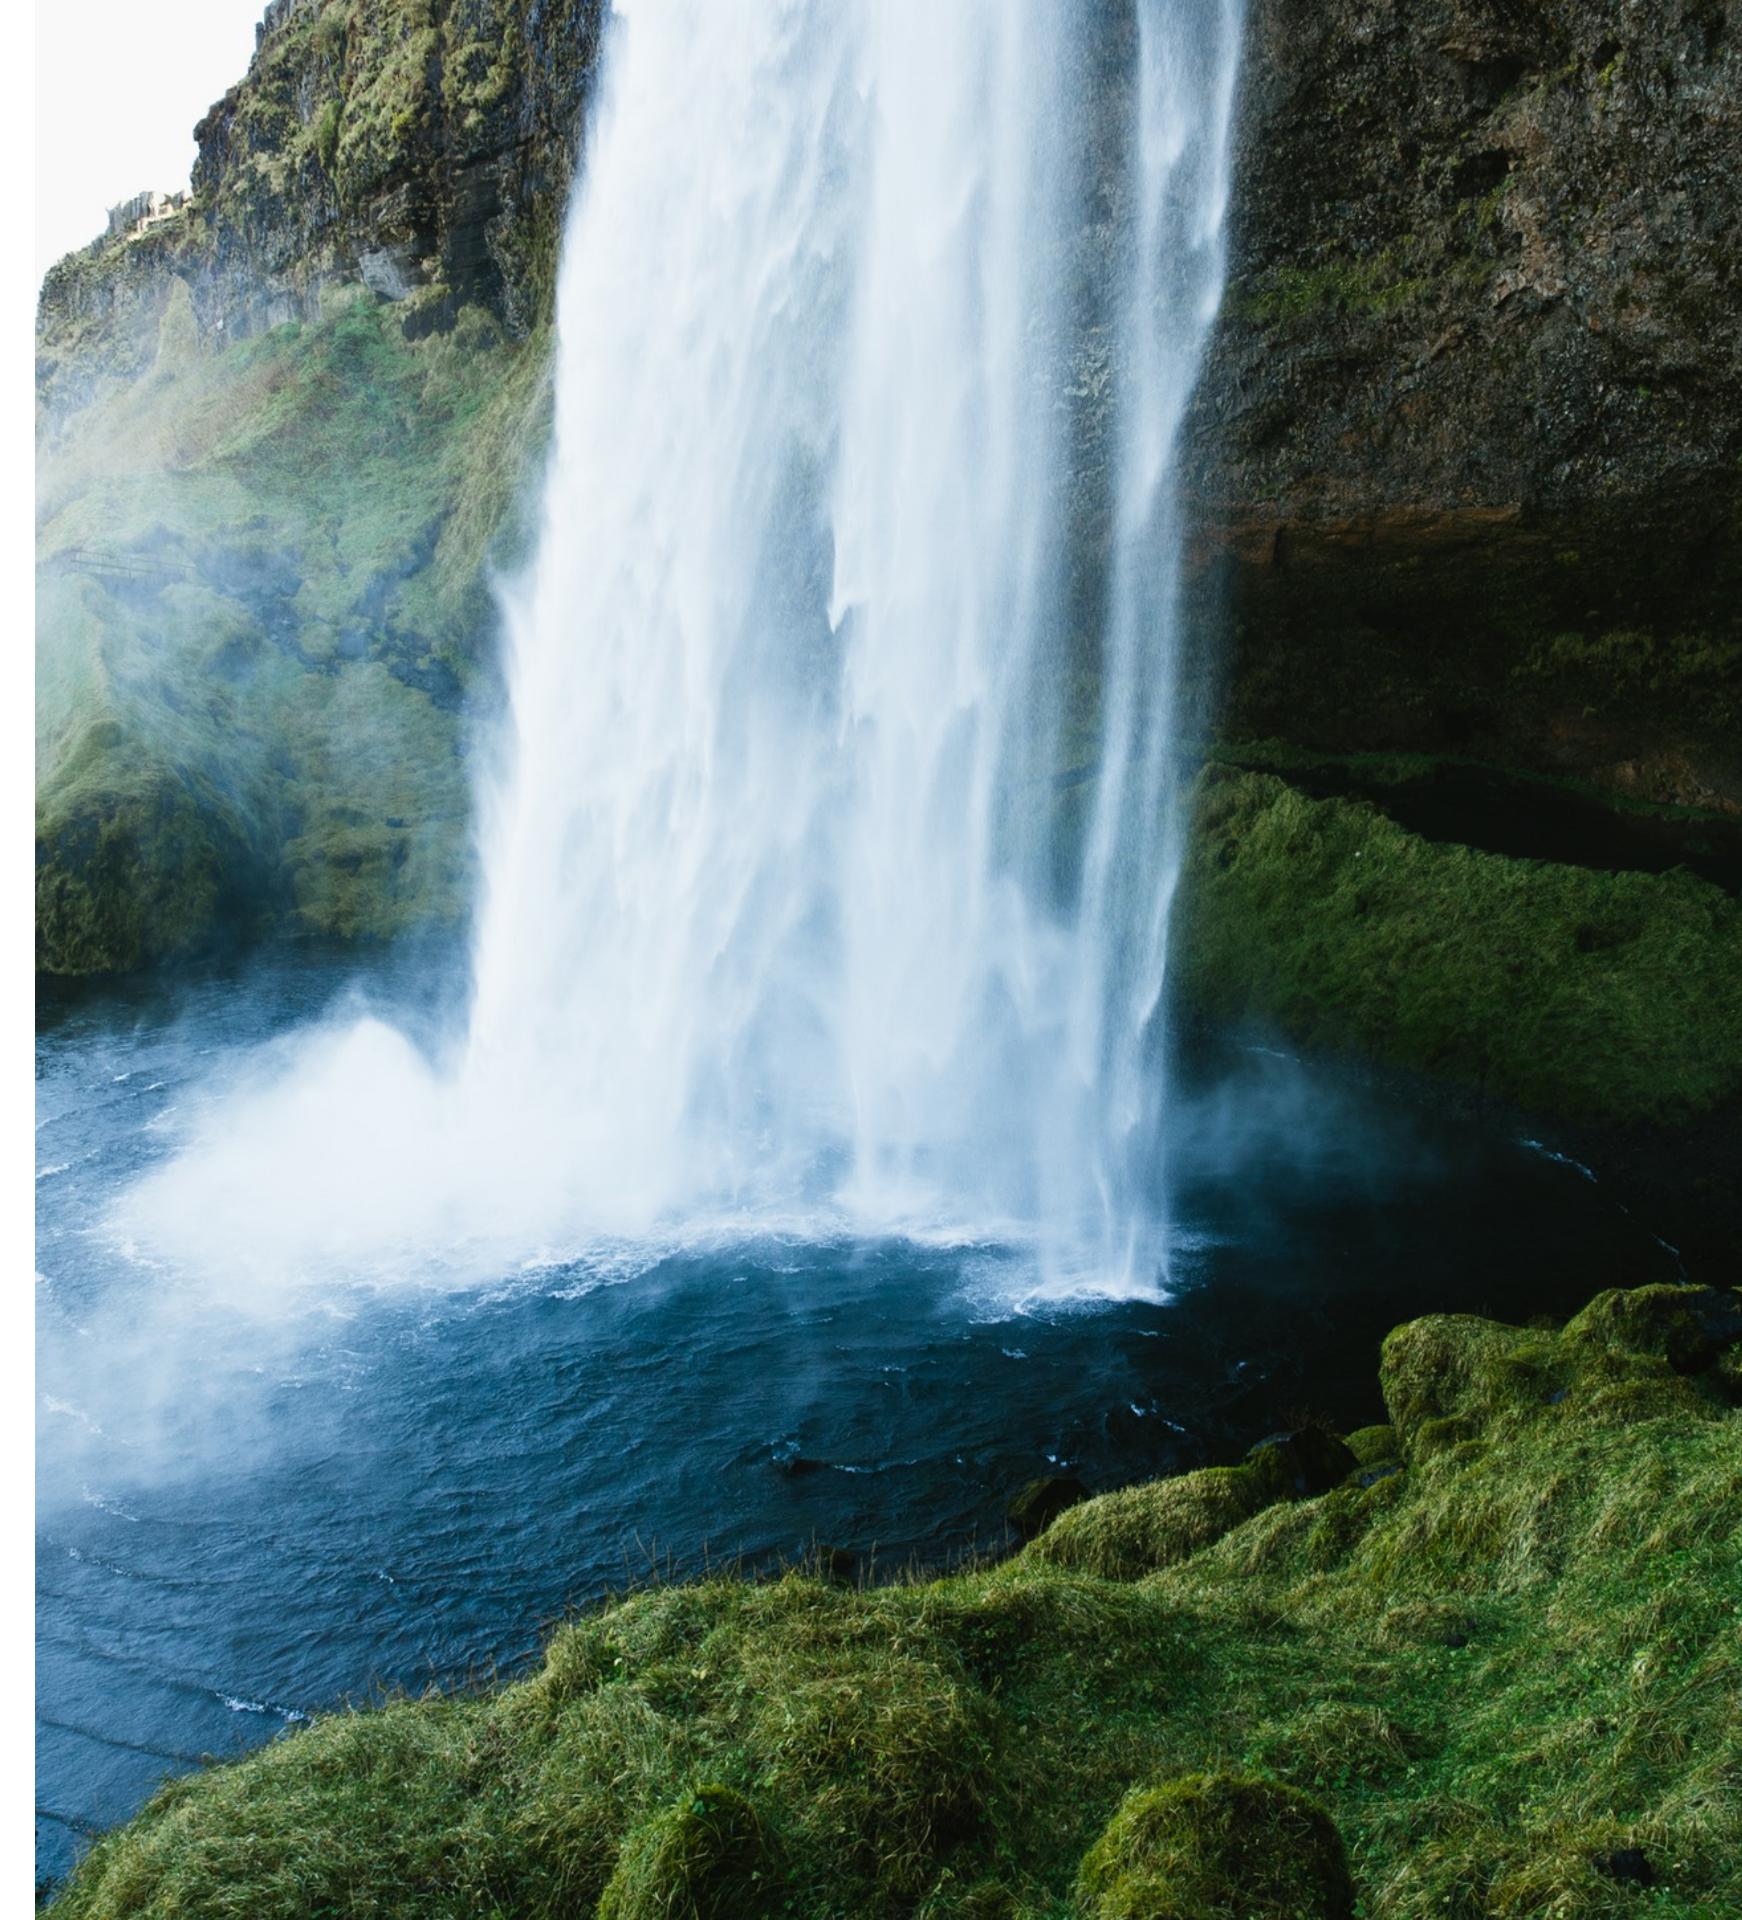


REACTIVE STREAMS

THE GLUE OF THE REAL-TIME ORGANIZATION

WHAT TO EXPECT

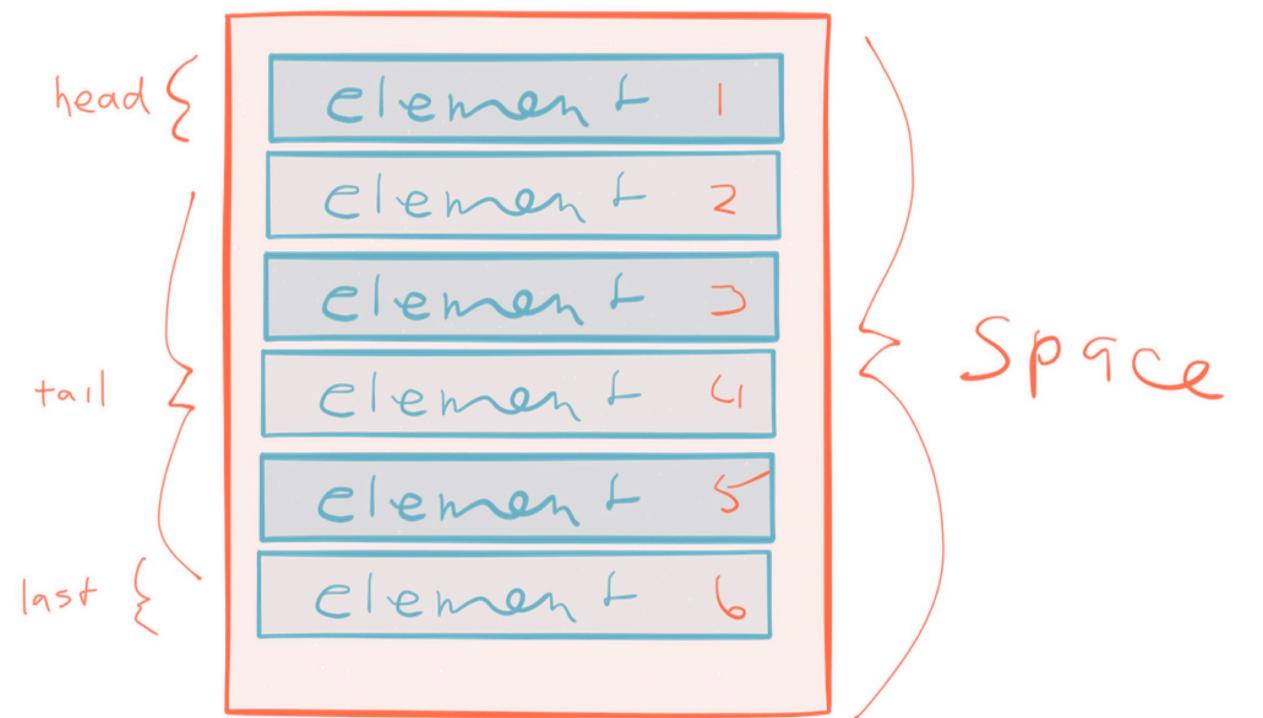
- » Introduction
- » A Journey into Reactive Streams
- » Diving into Akka Streams
- » Code walkthrough and demo
- » Q&A



**“IF ANALYTICS ARE THE
BRAIN OF AN
ORGANIZATION,
STREAMS ARE THE
HEART.”**

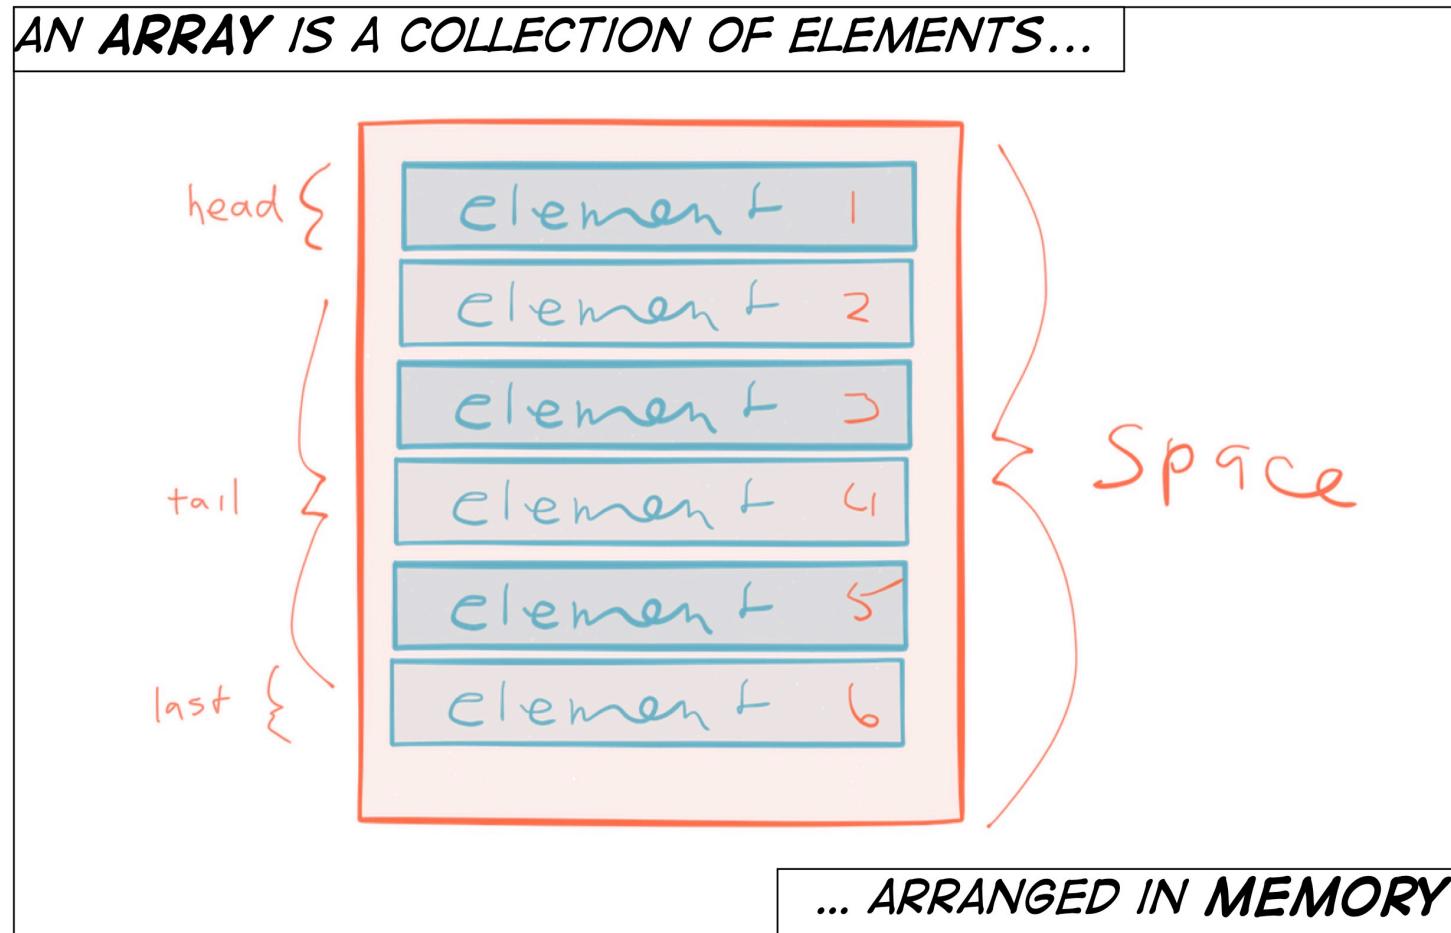
AN INTRODUCTION

AN ARRAY IS A COLLECTION OF ELEMENTS...

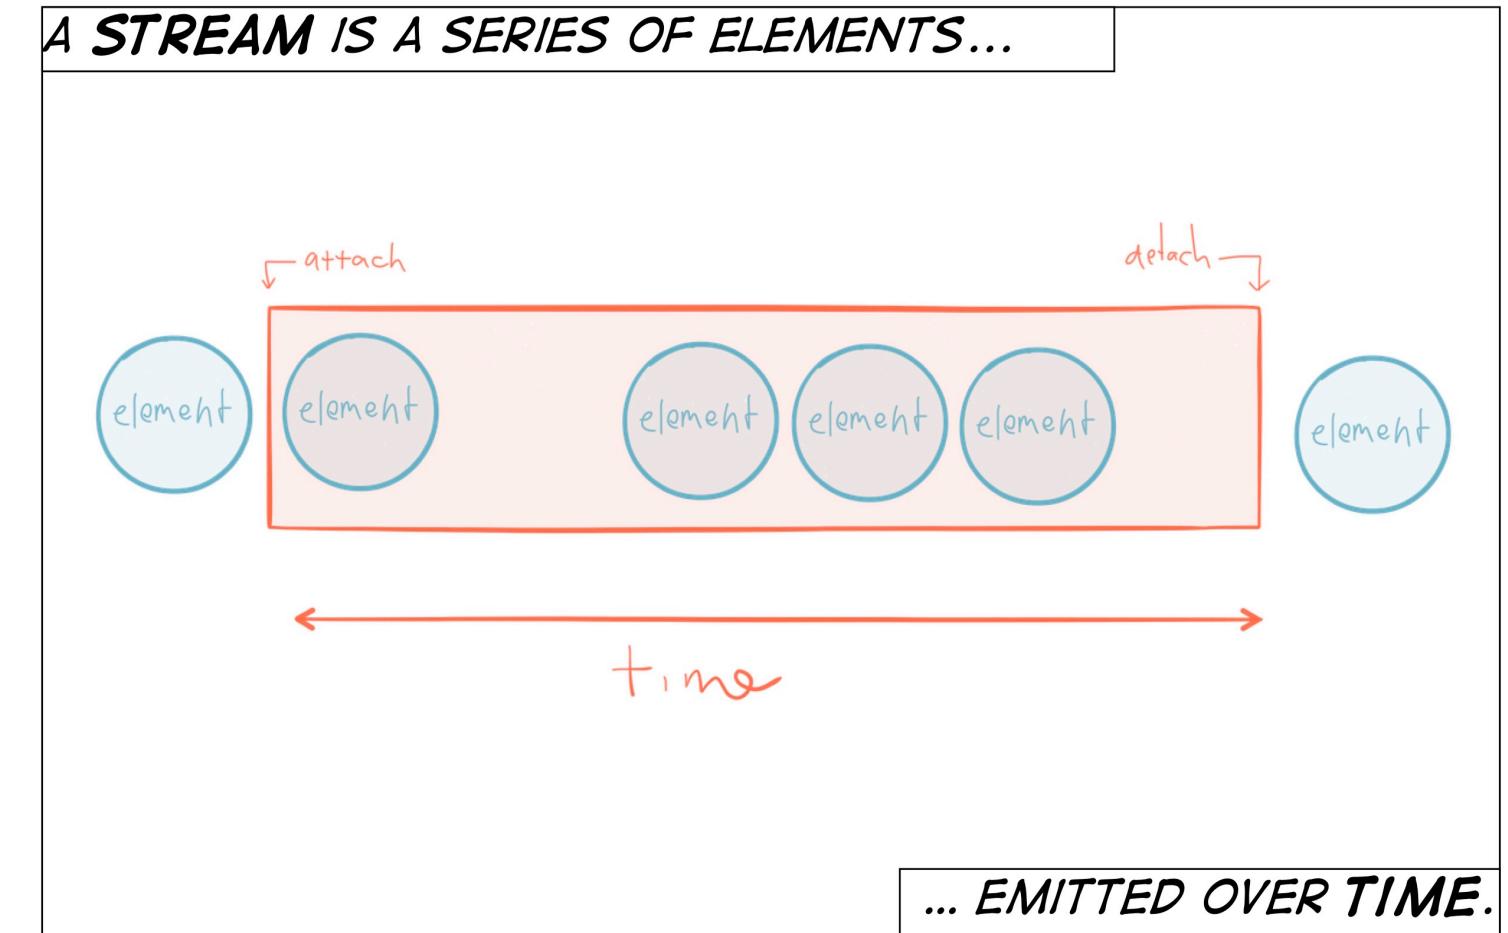


... ARRANGED IN MEMORY.

AN ARRAY IS A COLLECTION OF ELEMENTS...



A STREAM IS A SERIES OF ELEMENTS...



APPEAL OF STREAMS?

- » Per-event processing (\leq 1 second latencies)
- » Mini-batch processing (\leq 10 second latencies)
- » Batch processing (\leq 1 hour latencies)
- » Monitoring, analytics, complex event processing



CHALLENGES?

- » Ephemeral
- » Unbounded in size
- » Potential flooding
- » Unfamiliar

REACTIVE STREAMS

REACTIVE STREAMS?

- » Reactive Streams is a specification and low-level API for library developers

REACTIVE STREAMS?

- » Reactive Streams is a specification and low-level API for library developers
- » Started as an initiative in late 2013 between engineers at Netflix, Pivotal, and Lightbend

REACTIVE STREAMS?

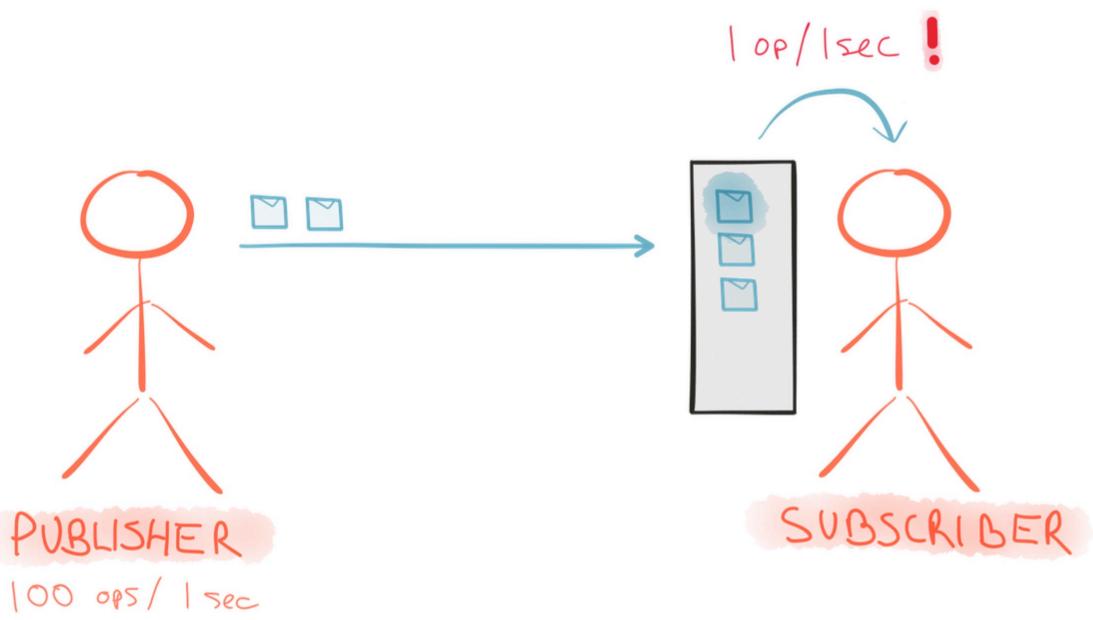
- » Reactive Streams is a specification and low-level API for library developers
- » Started as an initiative in late 2013 between engineers at Netflix, Pivotal, and Lightbend
- » Aimed to address two critical challenges
 - » Interoperability between streaming libraries
 - » Flow control over an asynchronous boundary

WHAT IS REACTIVE STREAMS?

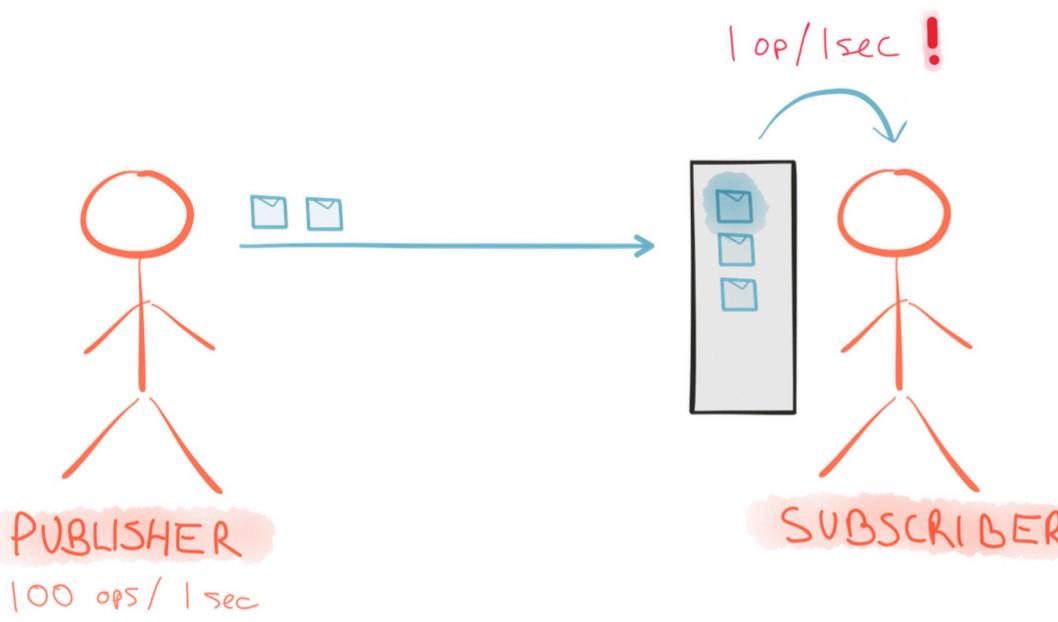
- » TCK (Technology Compatibility Kit)
- » API (JVM, JavaScript)
- » Specifications for library developers

FLOW CONTROL

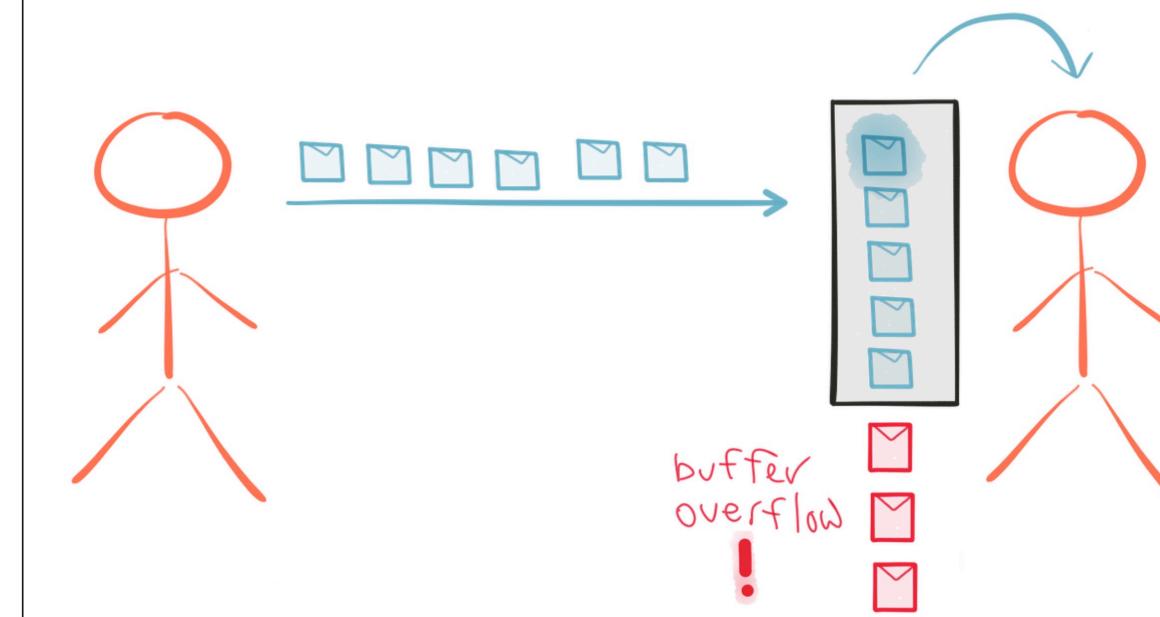
A SUBSCRIBER USUALLY HAS SOME TYPE OF BUFFER.



A SUBSCRIBER USUALLY HAS SOME TYPE OF BUFFER.

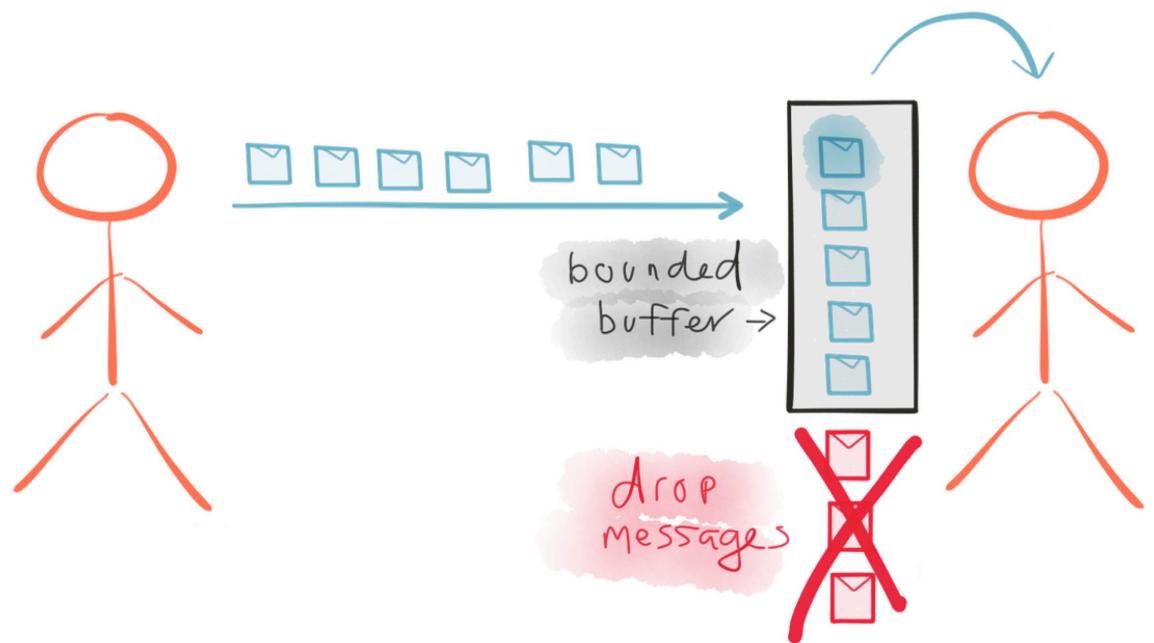


A FASTER PUBLISHER CAN EASILY OVERWHELM THE BUFFER OF A SLOW SUBSCRIBER...



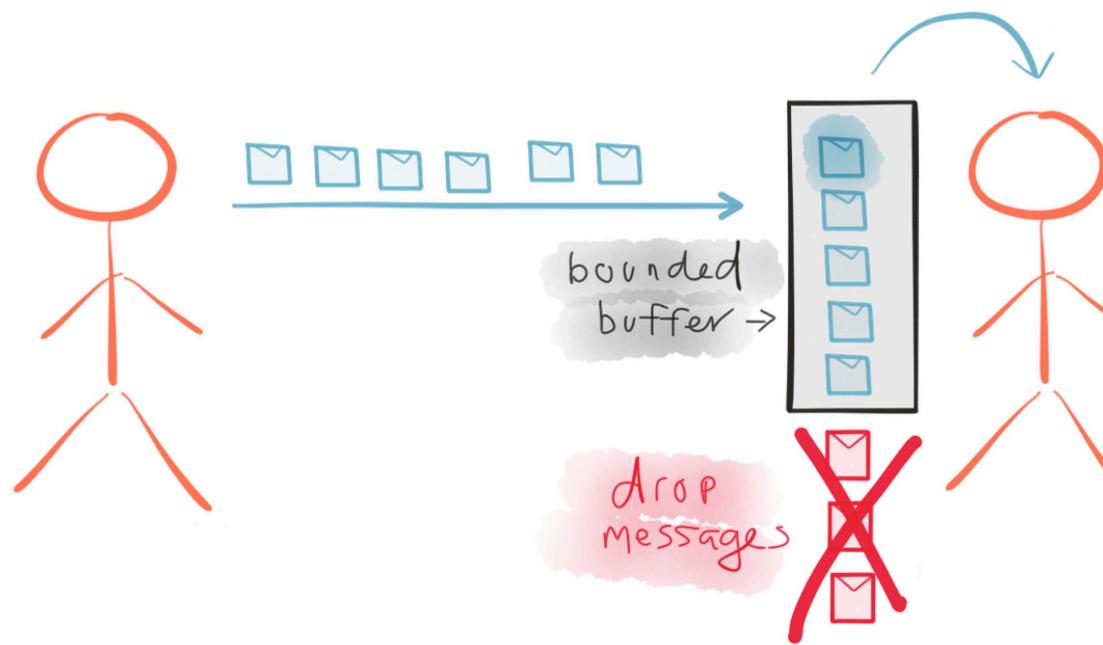
... WHICH CAN LEAD TO A CASCADING FAILURE THROUGH THE ENTIRE SYSTEM.

OPTION 1: DROP MESSAGES.



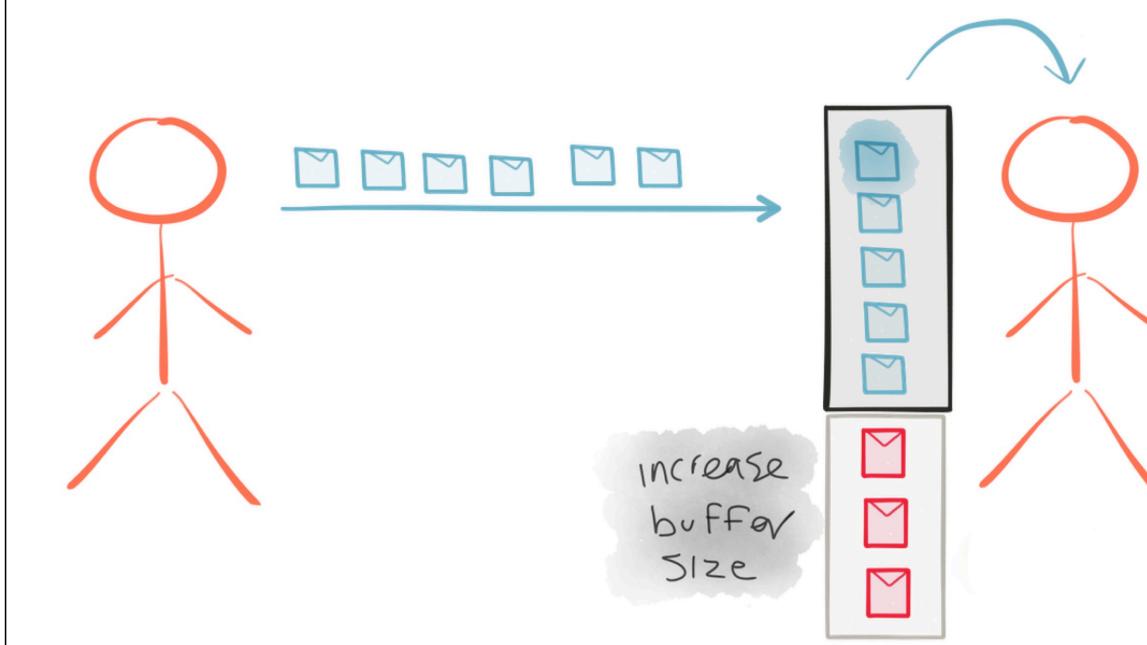
WORKS WELL FOR MONITORING AND OTHER
EPHEMERAL TASKS.

OPTION 1: DROP MESSAGES.



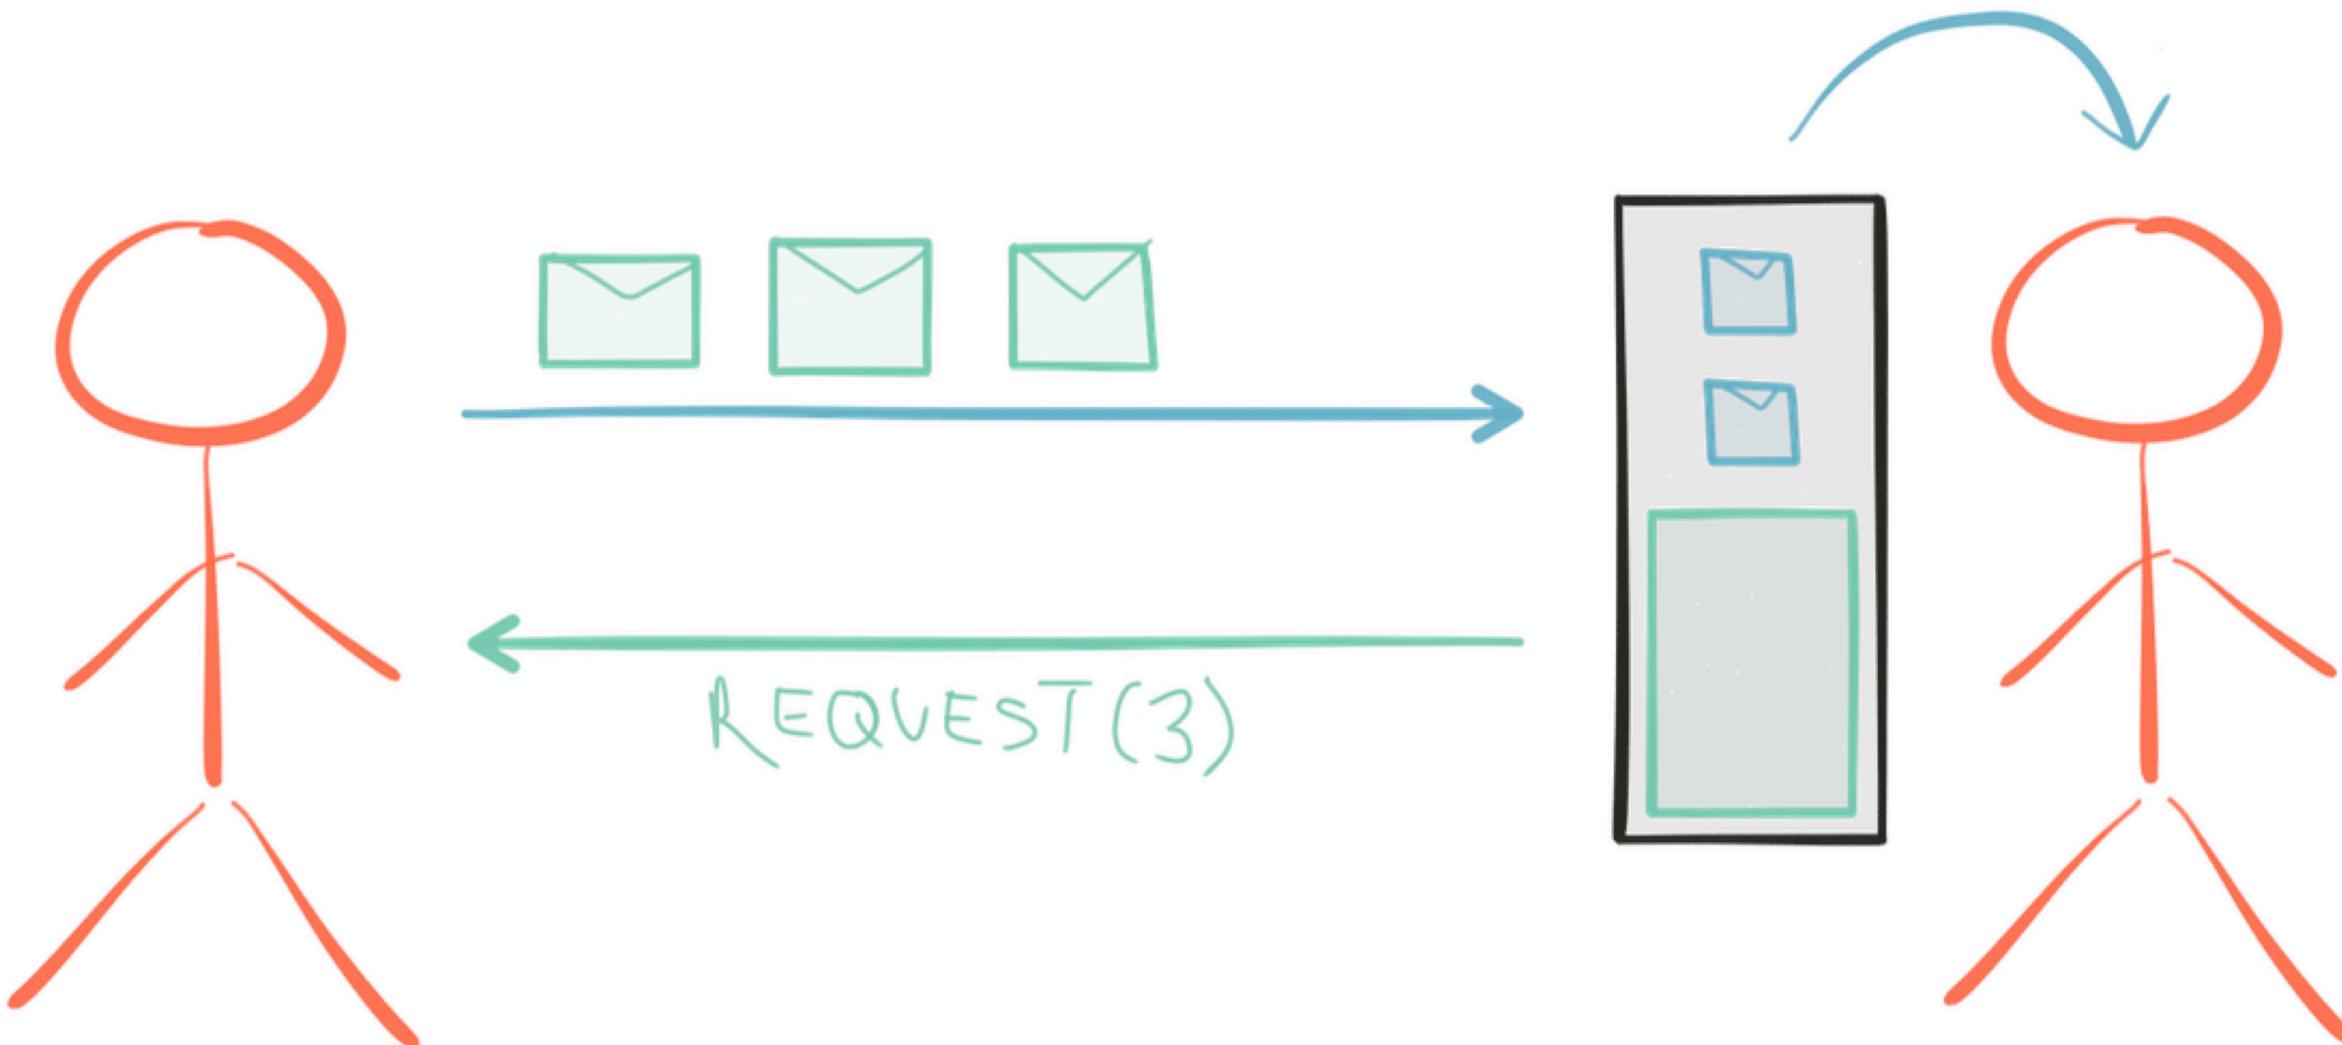
WORKS WELL FOR MONITORING AND OTHER
EPHEMERAL TASKS.

OPTION 2: INCREASE BUFFER SIZE.



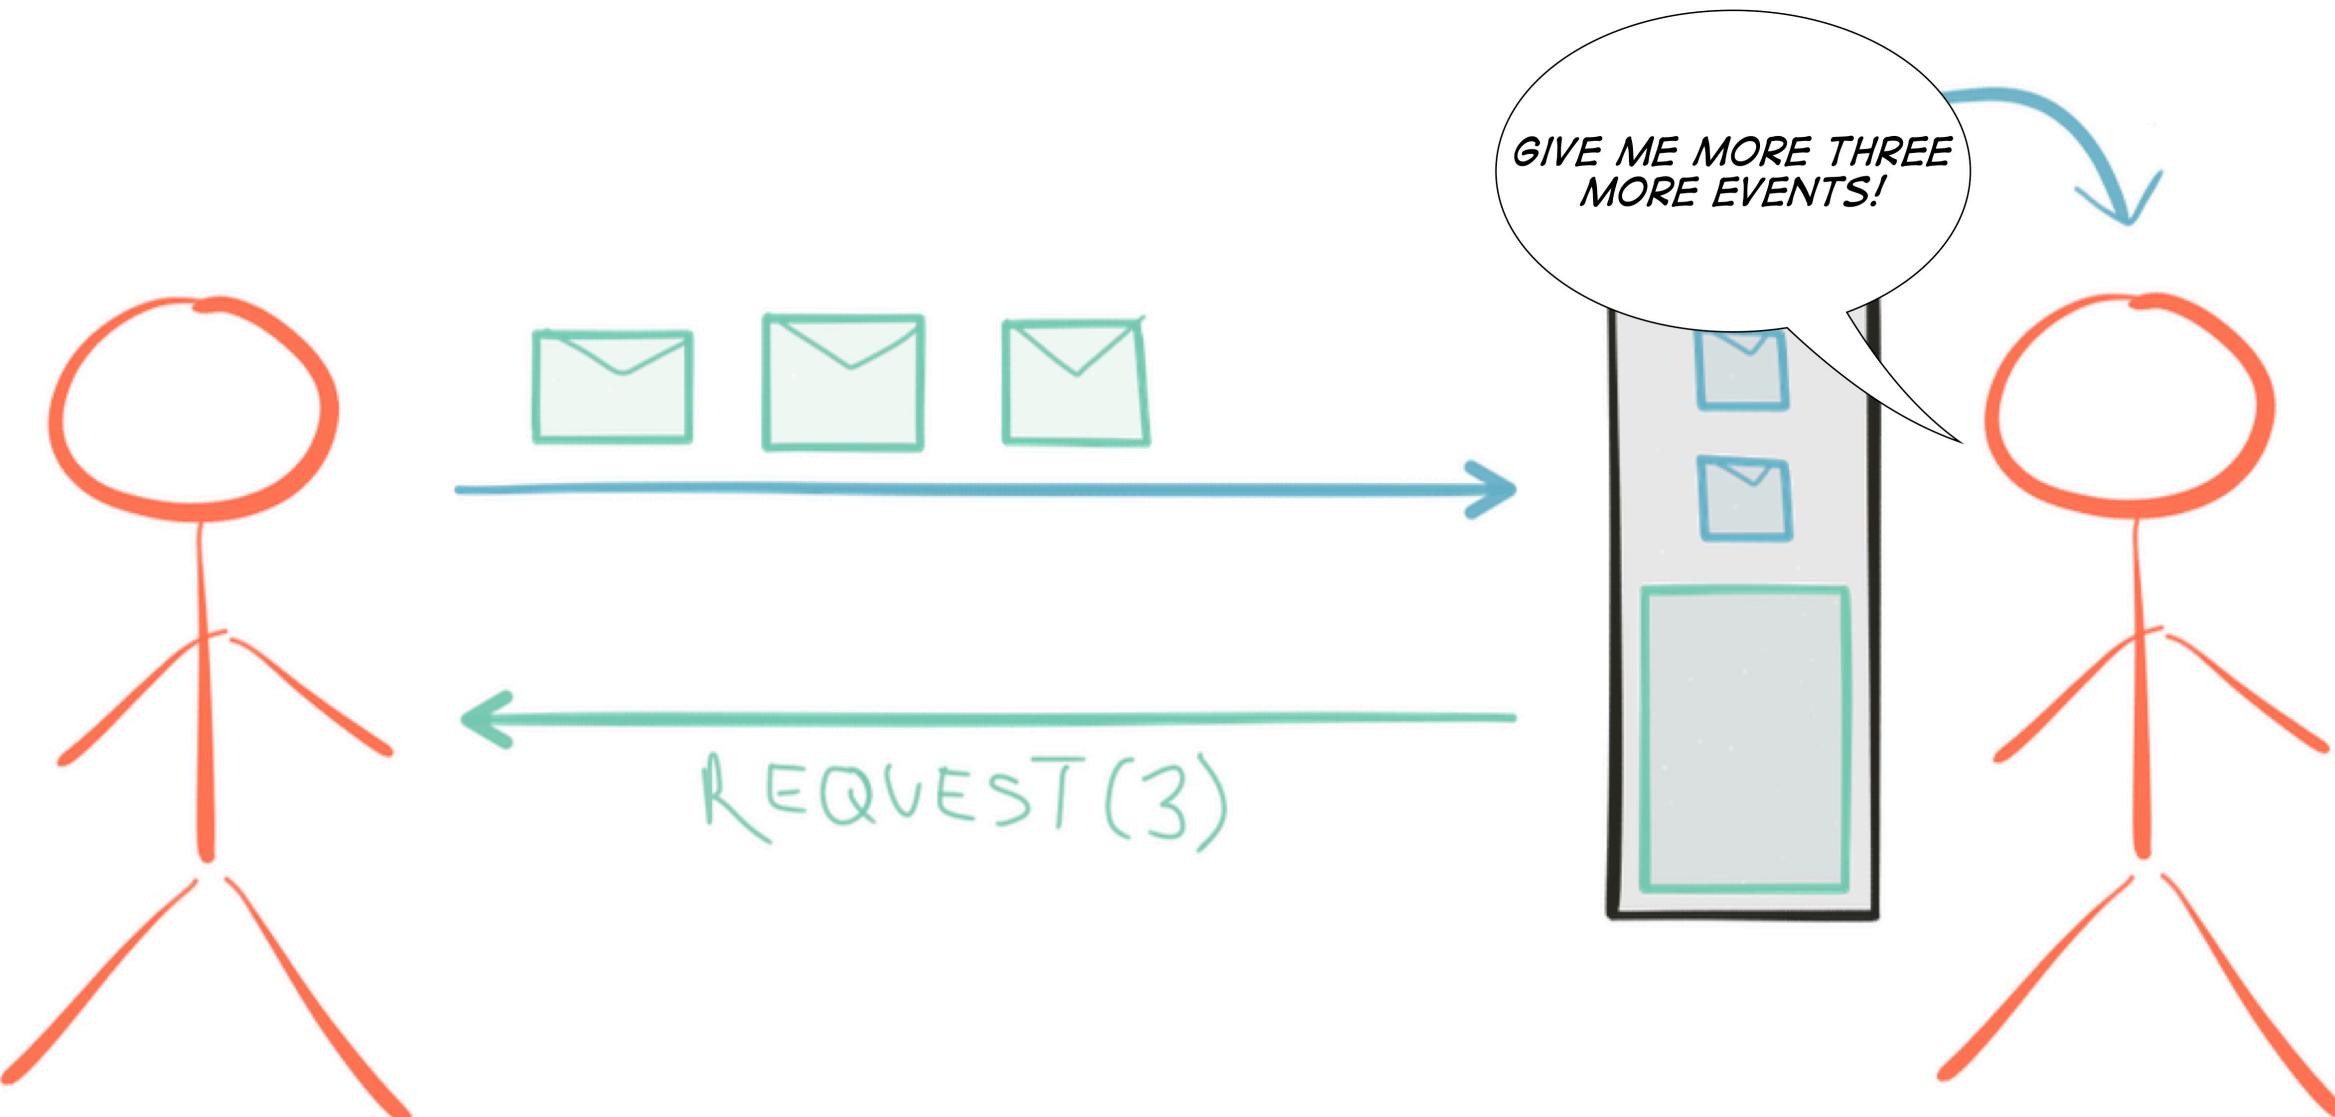
WORKS WELL UNTIL YOU RUN OUT OF
MEMORY!

IDEAL SOLUTION: PULL-BASED BACKPRESSURE.



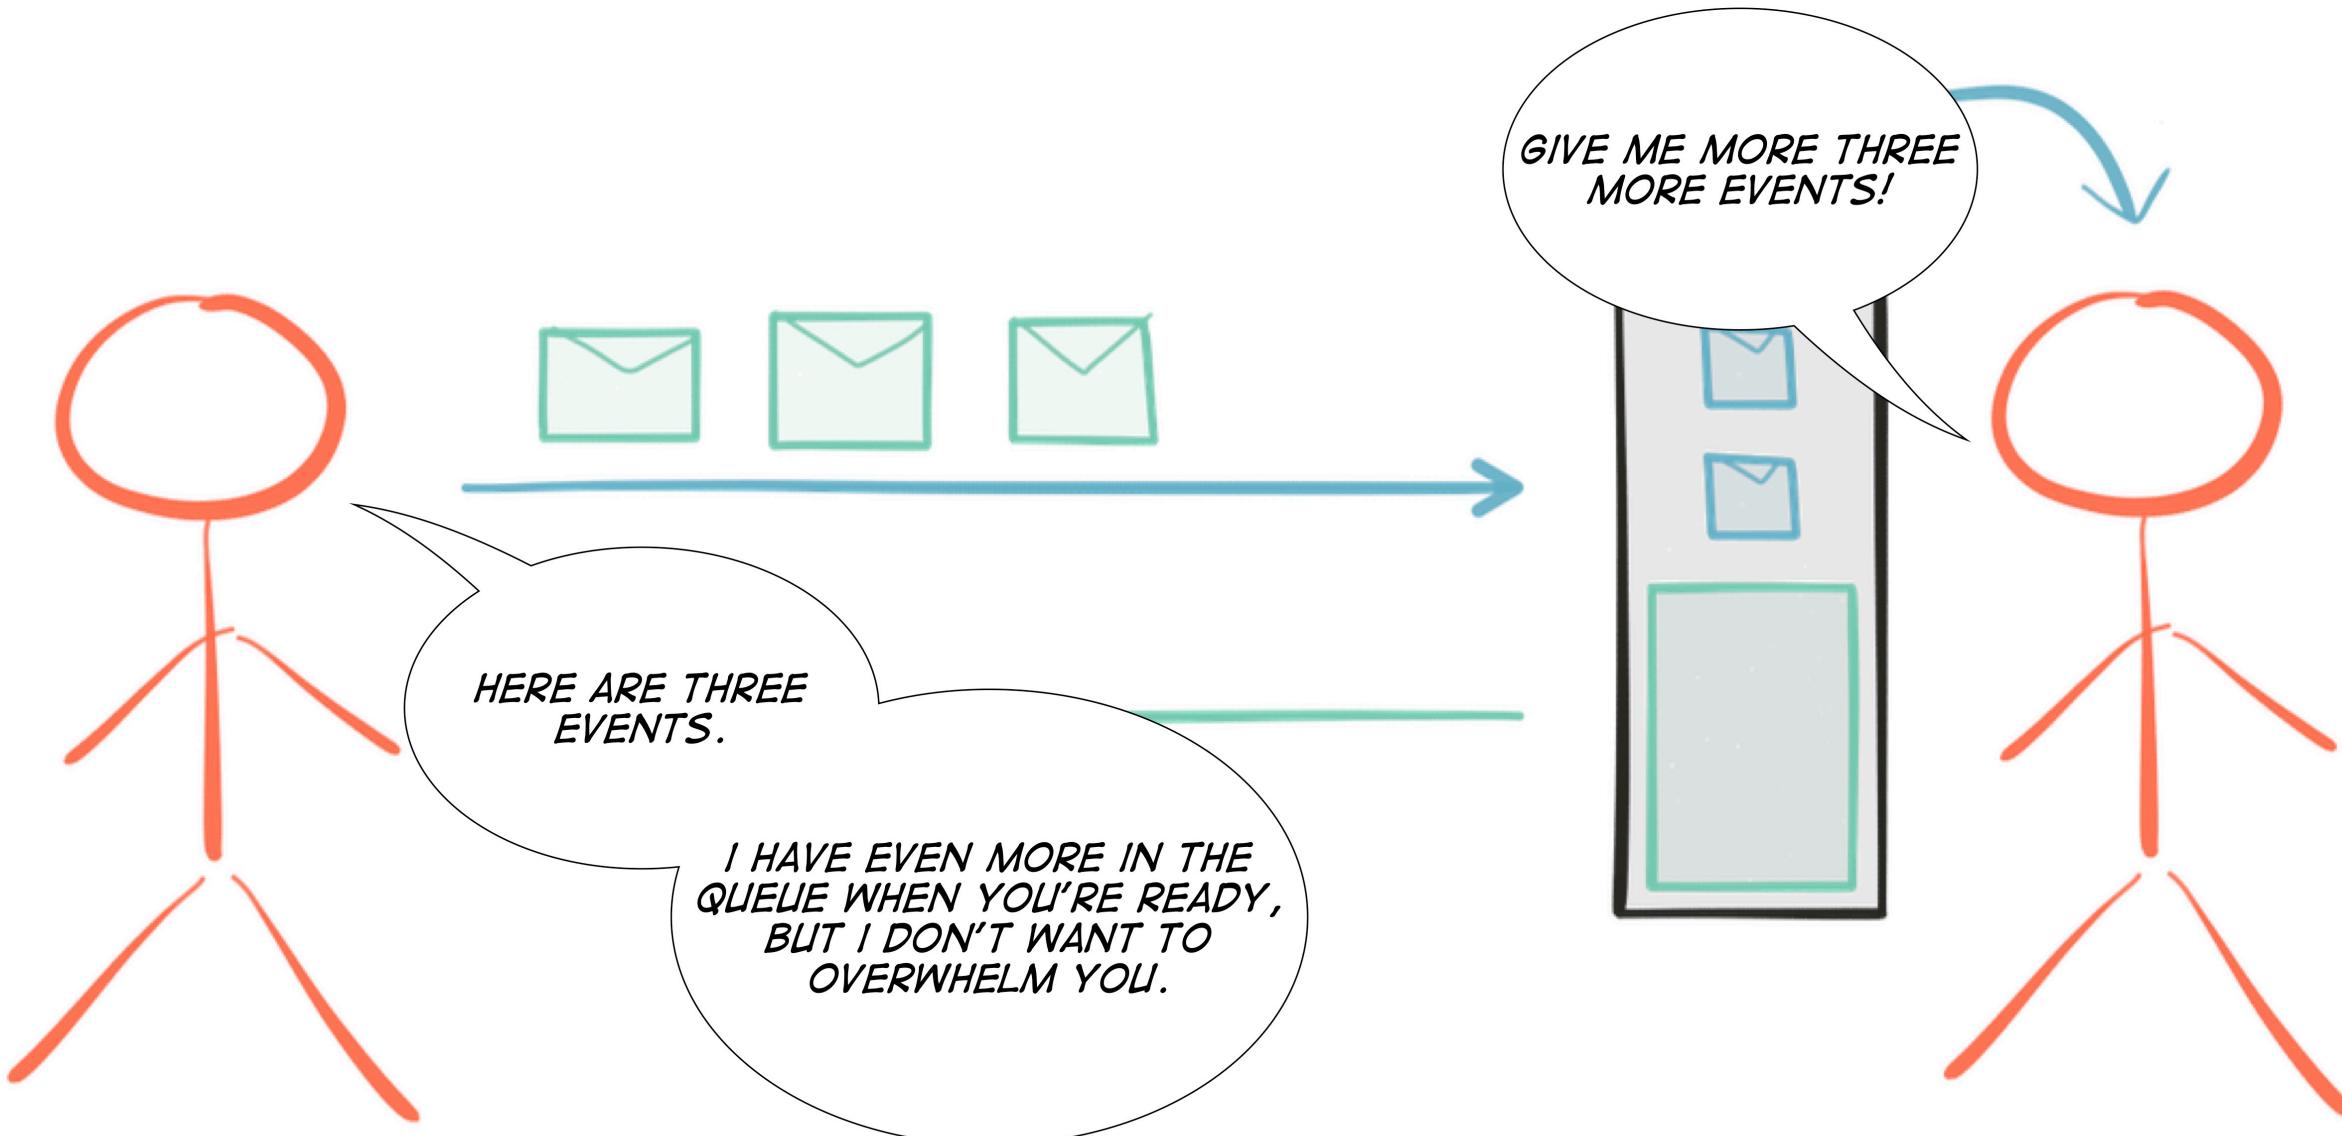
DEMAND FLOWS UPSTREAM, EVENTS FLOW DOWNSTREAM.

IDEAL SOLUTION: PULL-BASED BACKPRESSURE.



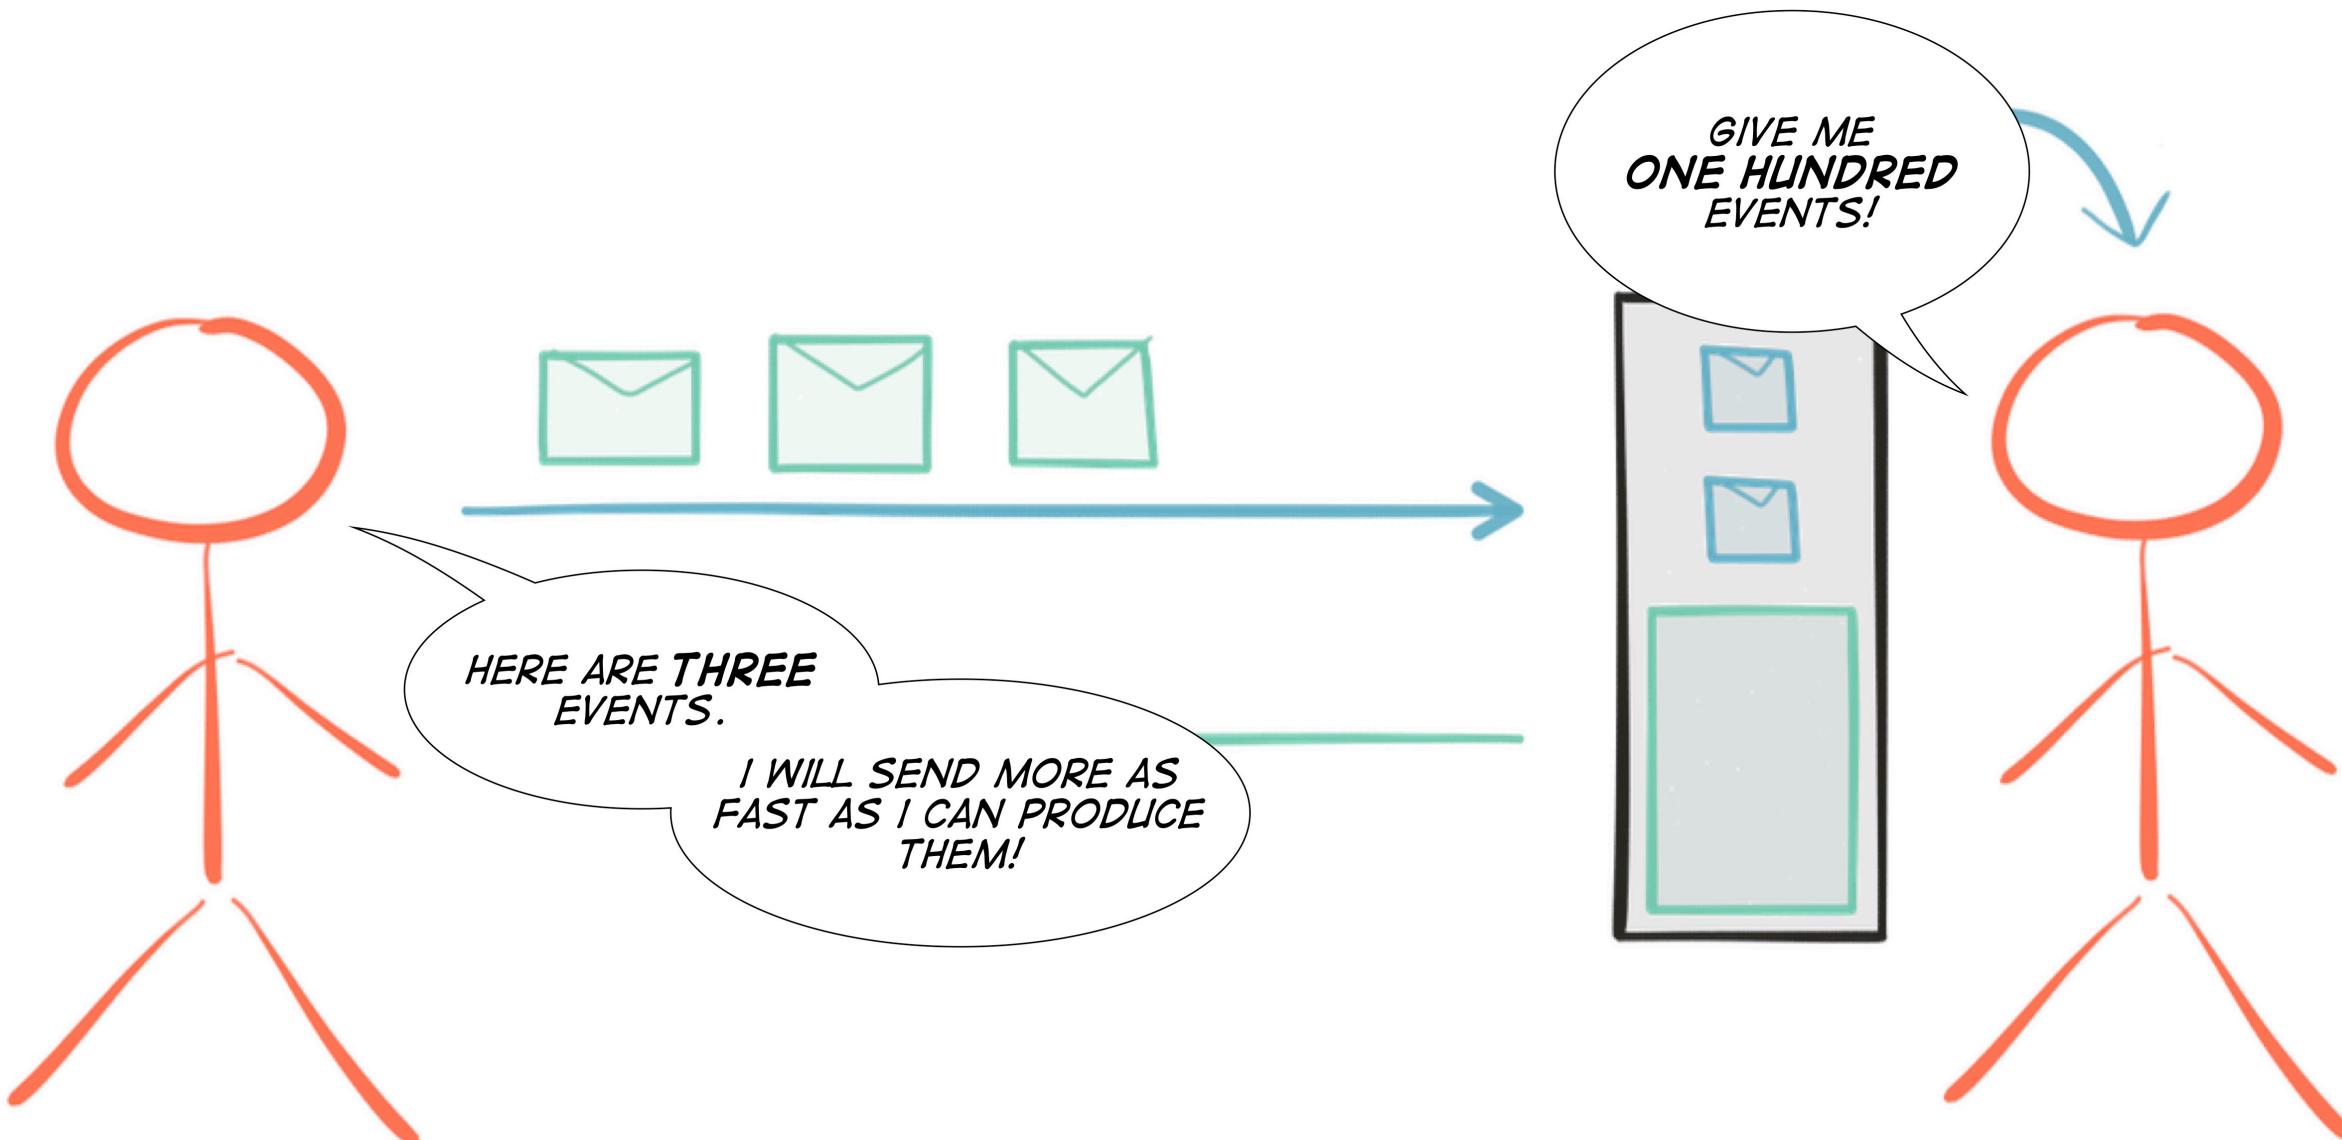
DEMAND FLOWS UPSTREAM, EVENTS FLOW DOWNSTREAM.

IDEAL SOLUTION: PULL-BASED BACKPRESSURE.



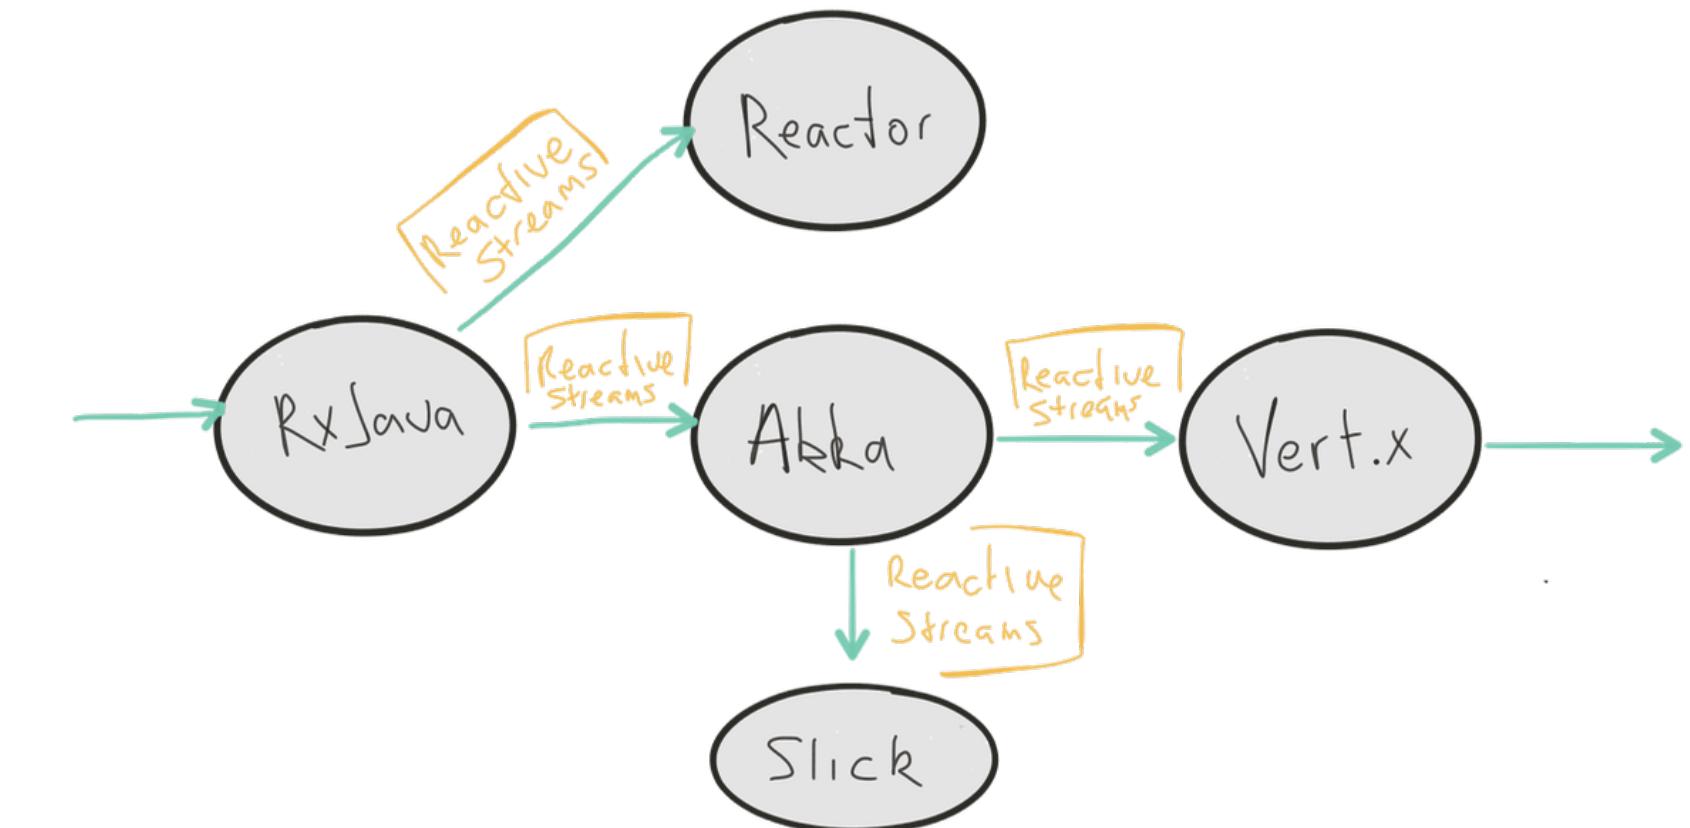
DEMAND FLOWS UPSTREAM, EVENTS FLOW DOWNSTREAM.

PULL-BASED BACKPRESSURE DYNAMICALLY CHANGES TO PUSH WHEN MORE DEMAND THAN SUPPLY.



INTEROP

- » RxJava (Netflix)
- » Reactor (Pivotal)
- » Vert.x (RedHat)
- » Akka Streams (Lightbend)
- » Slick (Lightbend)



REACTIVE STREAMS

Visit the Reactive Streams website for more information.

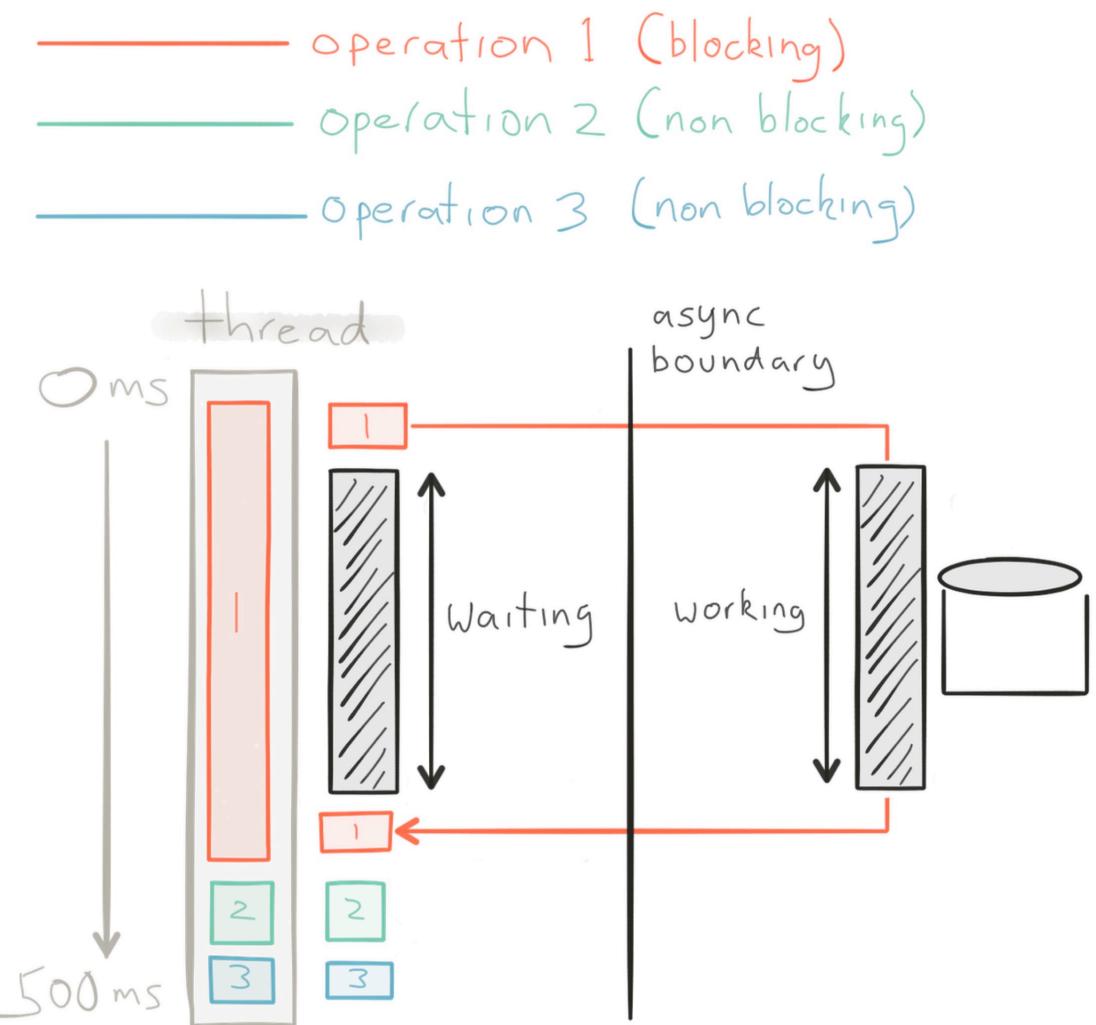
<http://www.reactive-streams.org/>

AKKA STREAMS

AKKA STREAMS

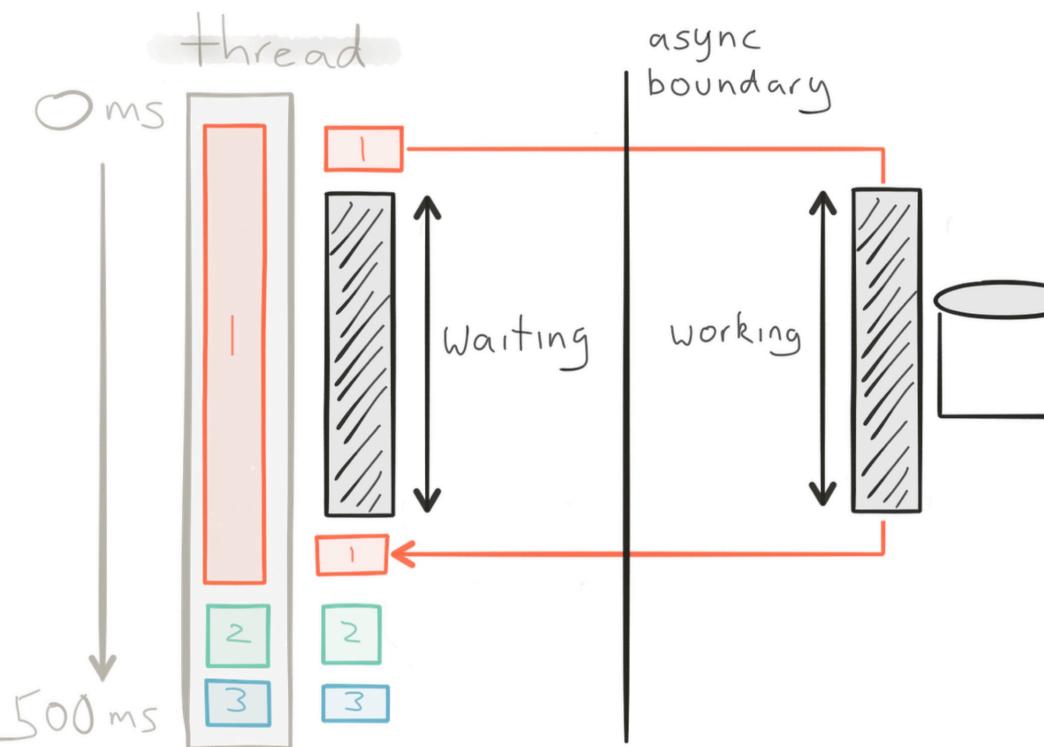
- » A library to express and run a chain of asynchronous processing steps acting on a sequence of elements
- » DSL for async/non-blocking stream processing
- » Backpressure enabled by default
- » Implements the Reactive Streams spec for interoperability
- » Scala and Java APIs

ASYNCHRONOUS BOUNDARY



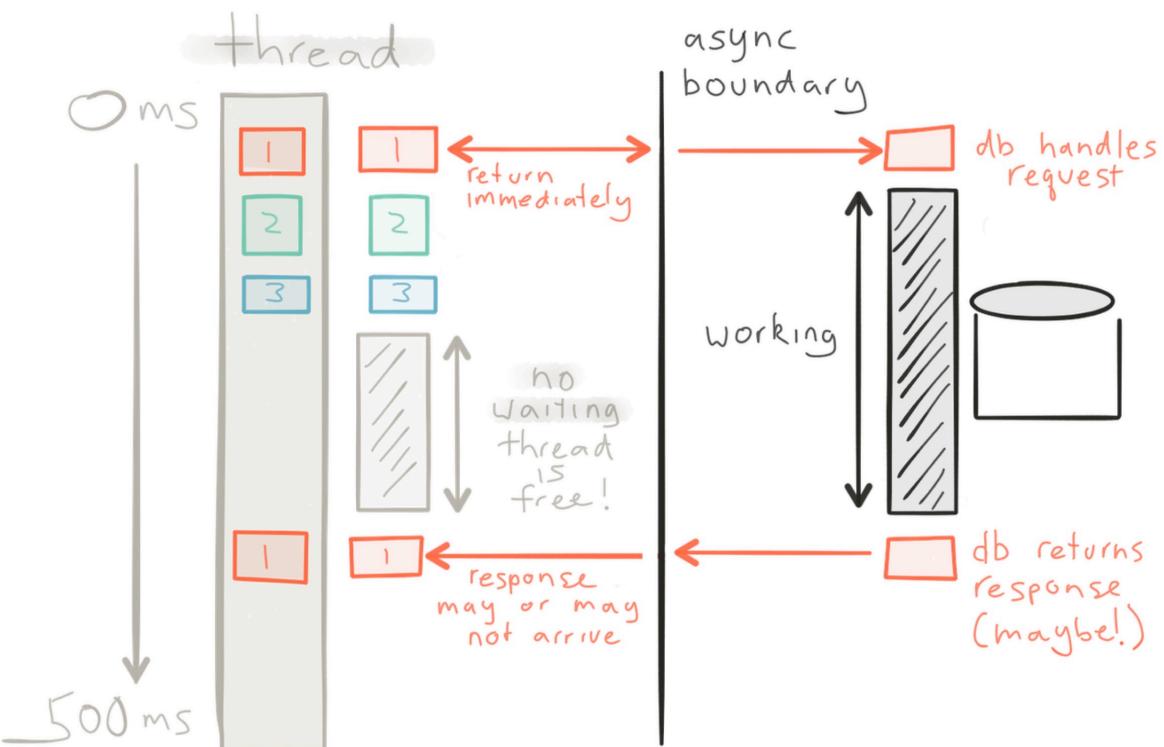
SYNCHRONOUS.

— operation 1 (blocking)
— operation 2 (non blocking)
— operation 3 (non blocking)



SYNCHRONOUS.

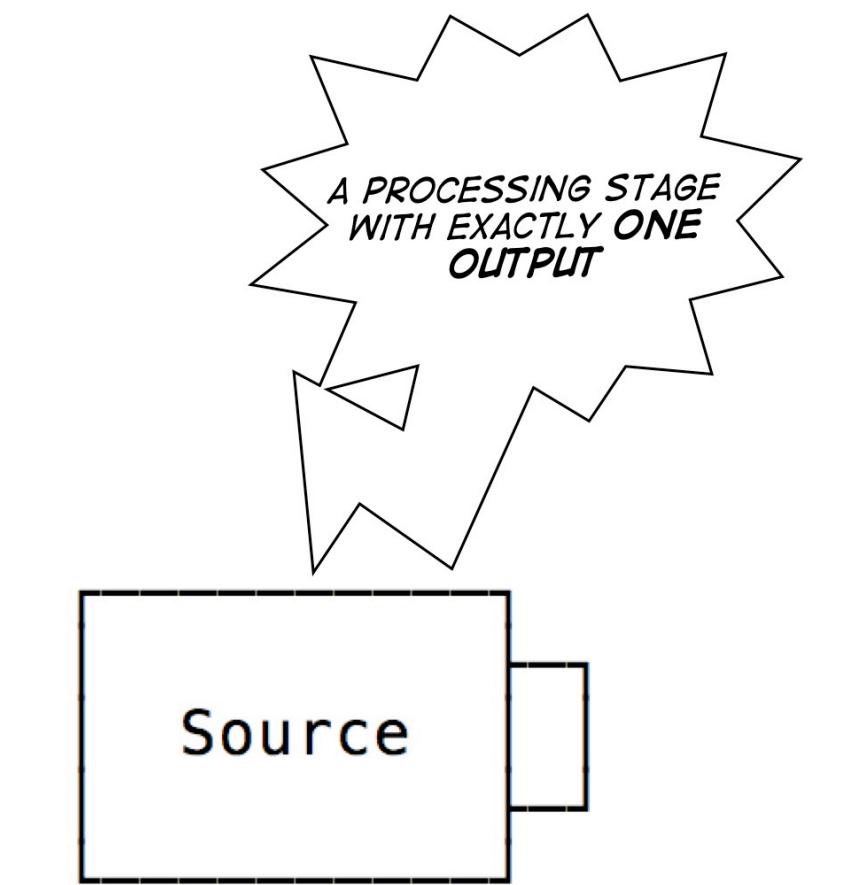
— operation 1 (non blocking)
— operation 2 (non blocking)
— operation 3 (non blocking)

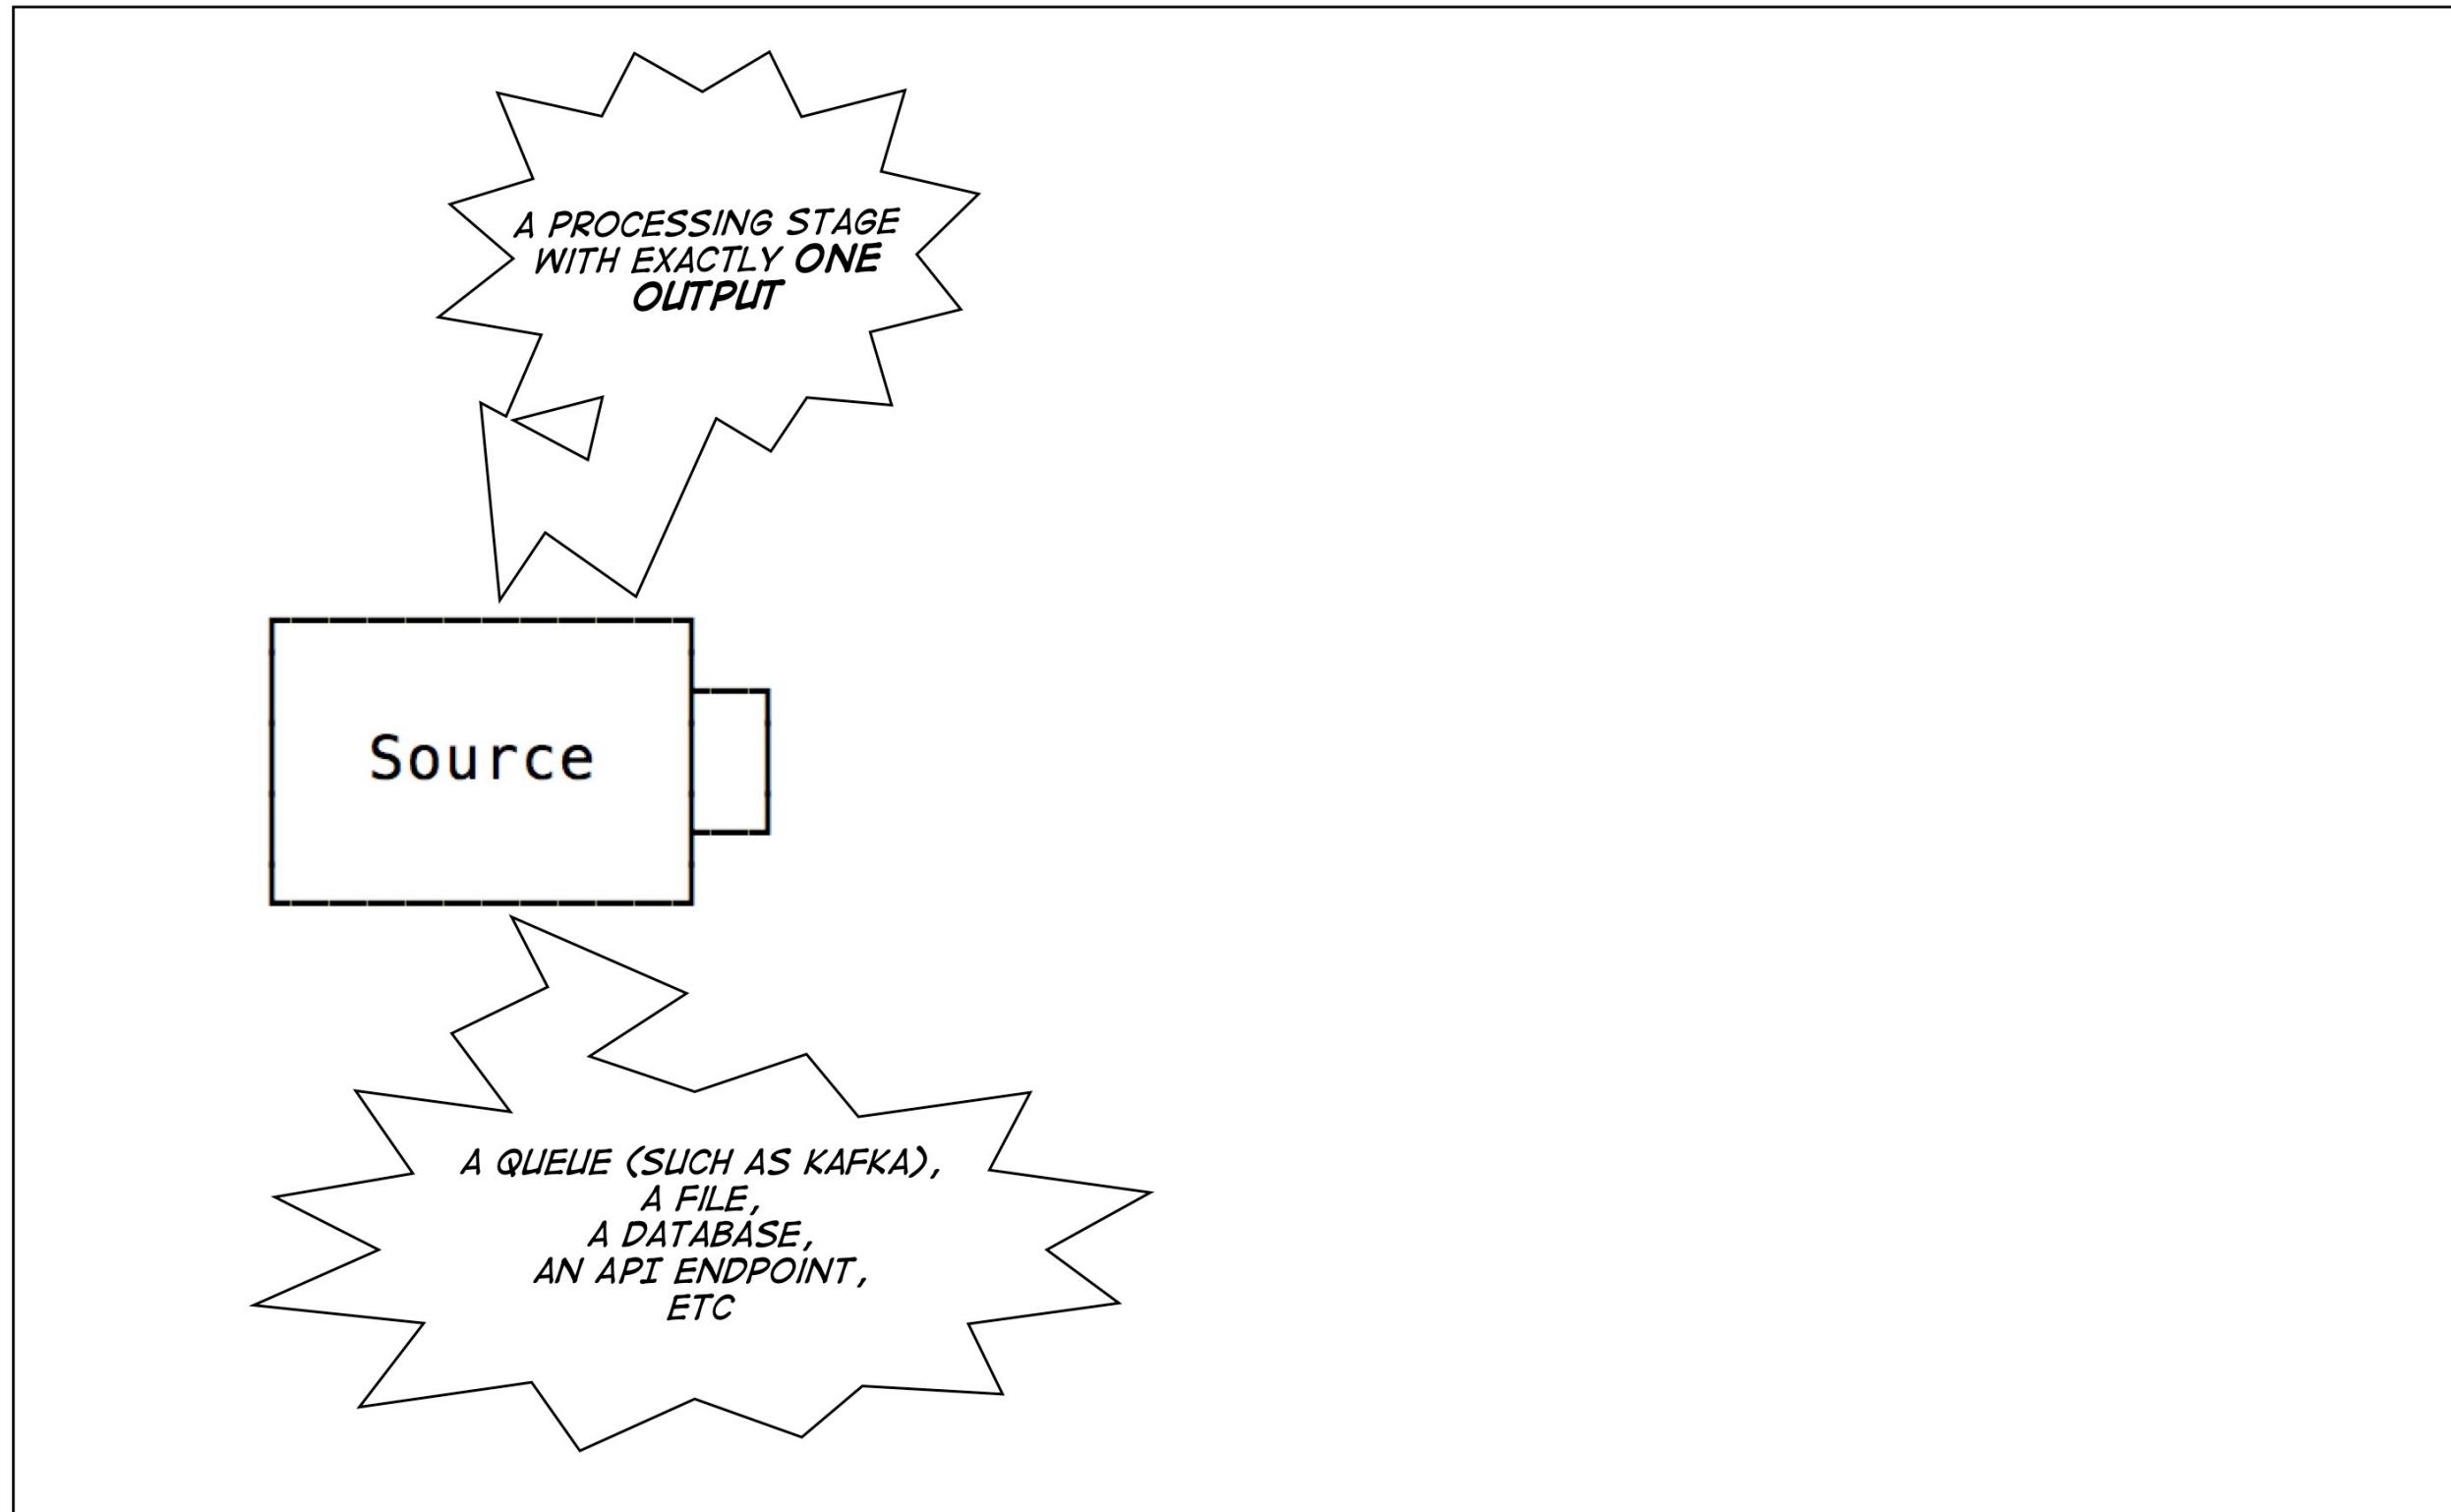


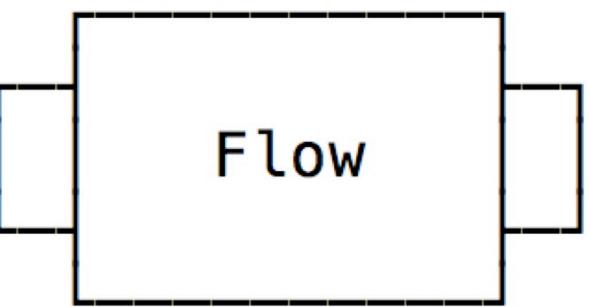
ASYNCHRONOUS.

BASICS

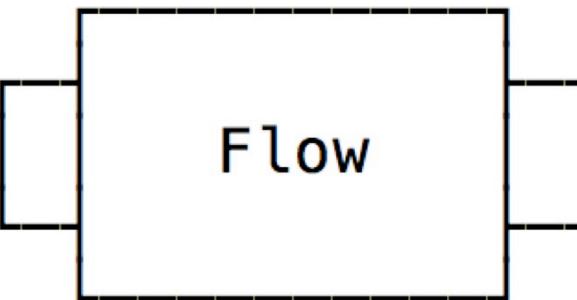




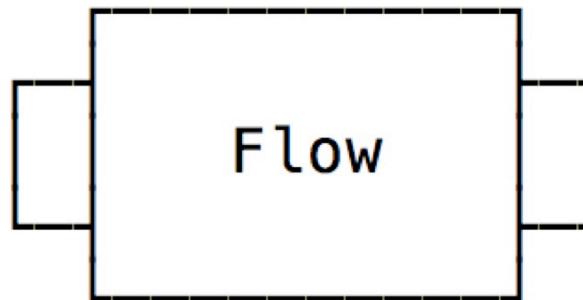




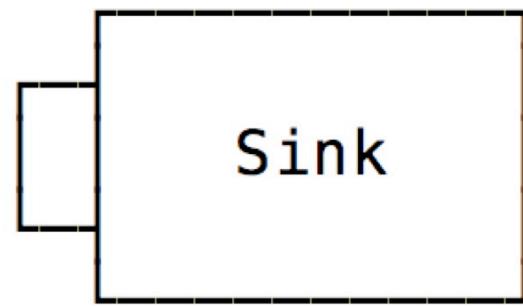
A PROCESSING STAGE
WITH EXACTLY
ONE INPUT
AND
ONE OUTPUT



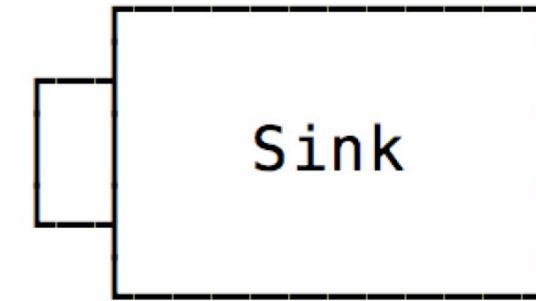
A PROCESSING STAGE
WITH EXACTLY
ONE INPUT
AND
ONE OUTPUT

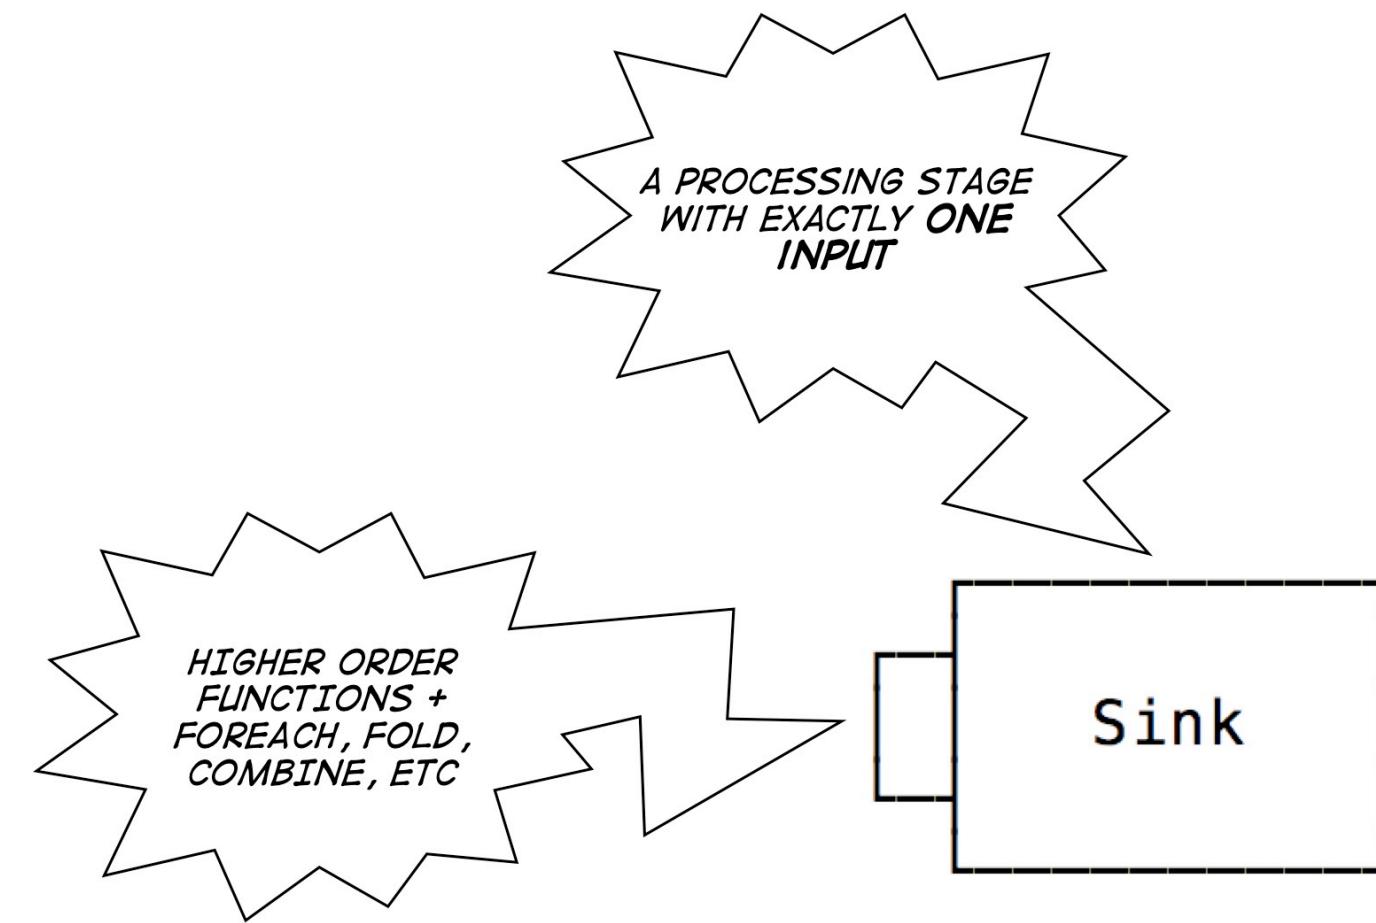


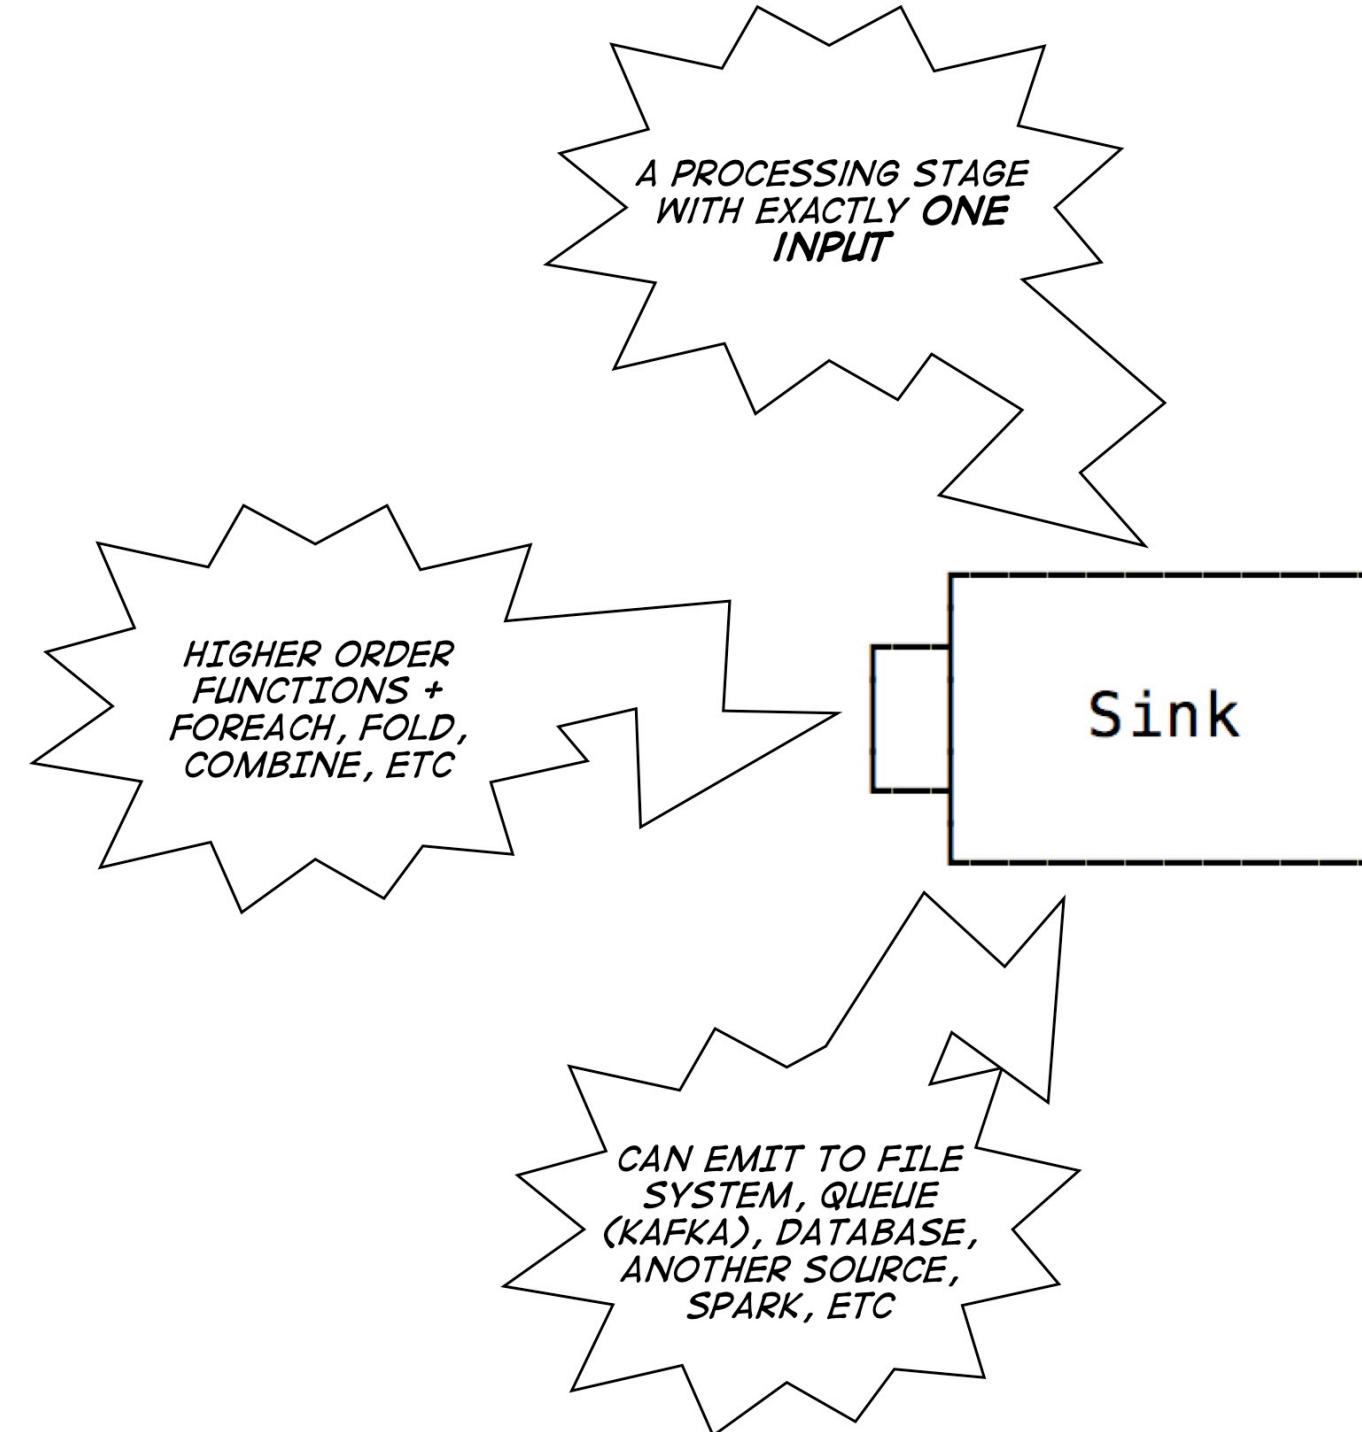
USES HIGHER ORDER
FUNCTIONS TO CONVERT
FROM AN INPUT TYPE TO AN
OUTPUT TYPE



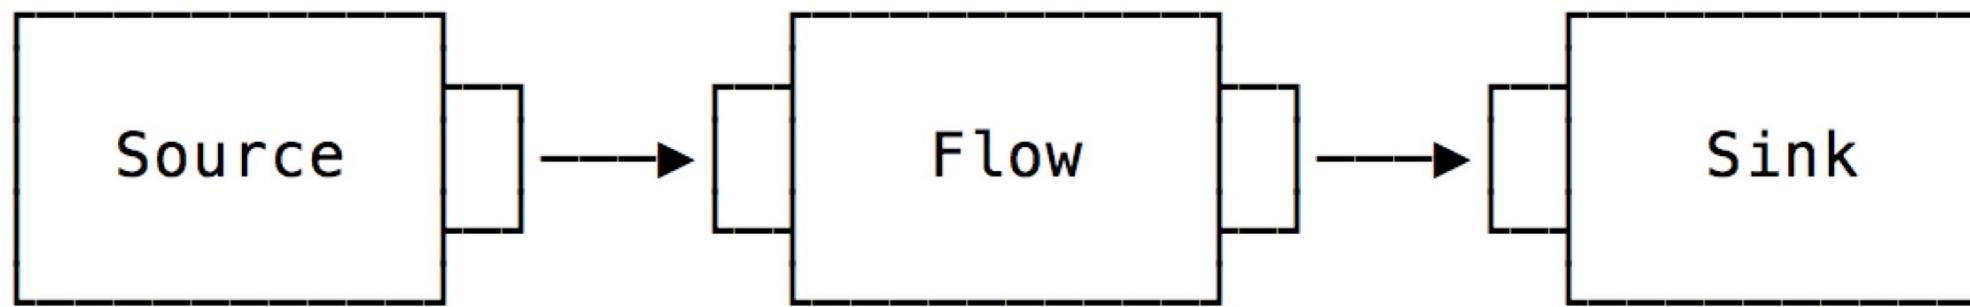
*A PROCESSING STAGE
WITH EXACTLY ONE
INPUT*





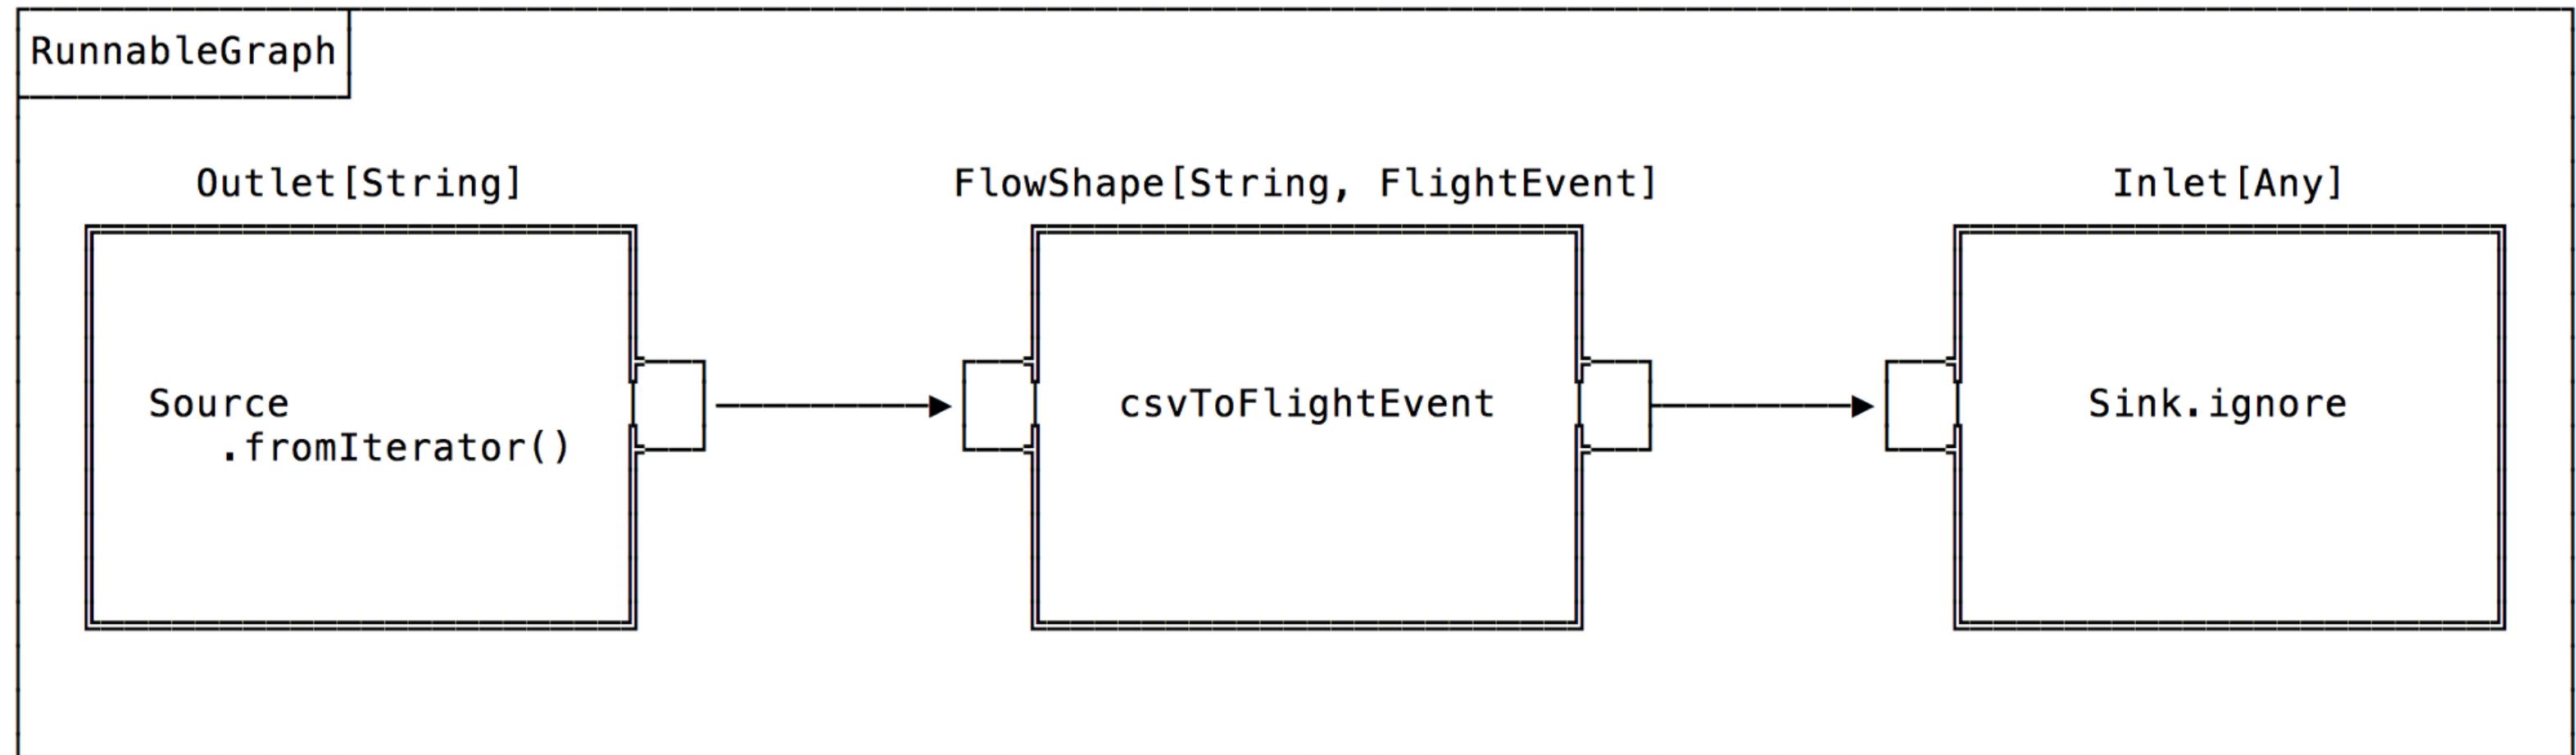


RUNNABLEGRAPH - ALL STAGES ARE CONNECTED TO A SOURCE AND A SINK.



LET'S ANALYZE FLIGHT DATA

1. Read in all flight data from 2008
2. Transform CSV rows to data structures
3. Compute the average delay per airline
4. Emit the results to console



```
val g: RunnableGraph[_] = RunnableGraph.fromGraph(GraphDSL.create() {
    implicit builder =>

    // Source
    val A: Outlet[String] = builder.add(Source.fromIterator(() => flightDelayLines)).out
    val B: FlowShape[String, FlightEvent] = builder.add(csvToFlightEvent)
    val C: Inlet[Any] = builder.add(Sink.ignore).in

    import GraphDSL.Implicits._ // allows us to build our graph using ~> combinators

    // Graph
    A ~> B ~> C

    ClosedShape // defines this as a "closed" graph, not exposing any inlets or outlets
})

g.run() // materializes and executes the blueprint
```

API DESIGN

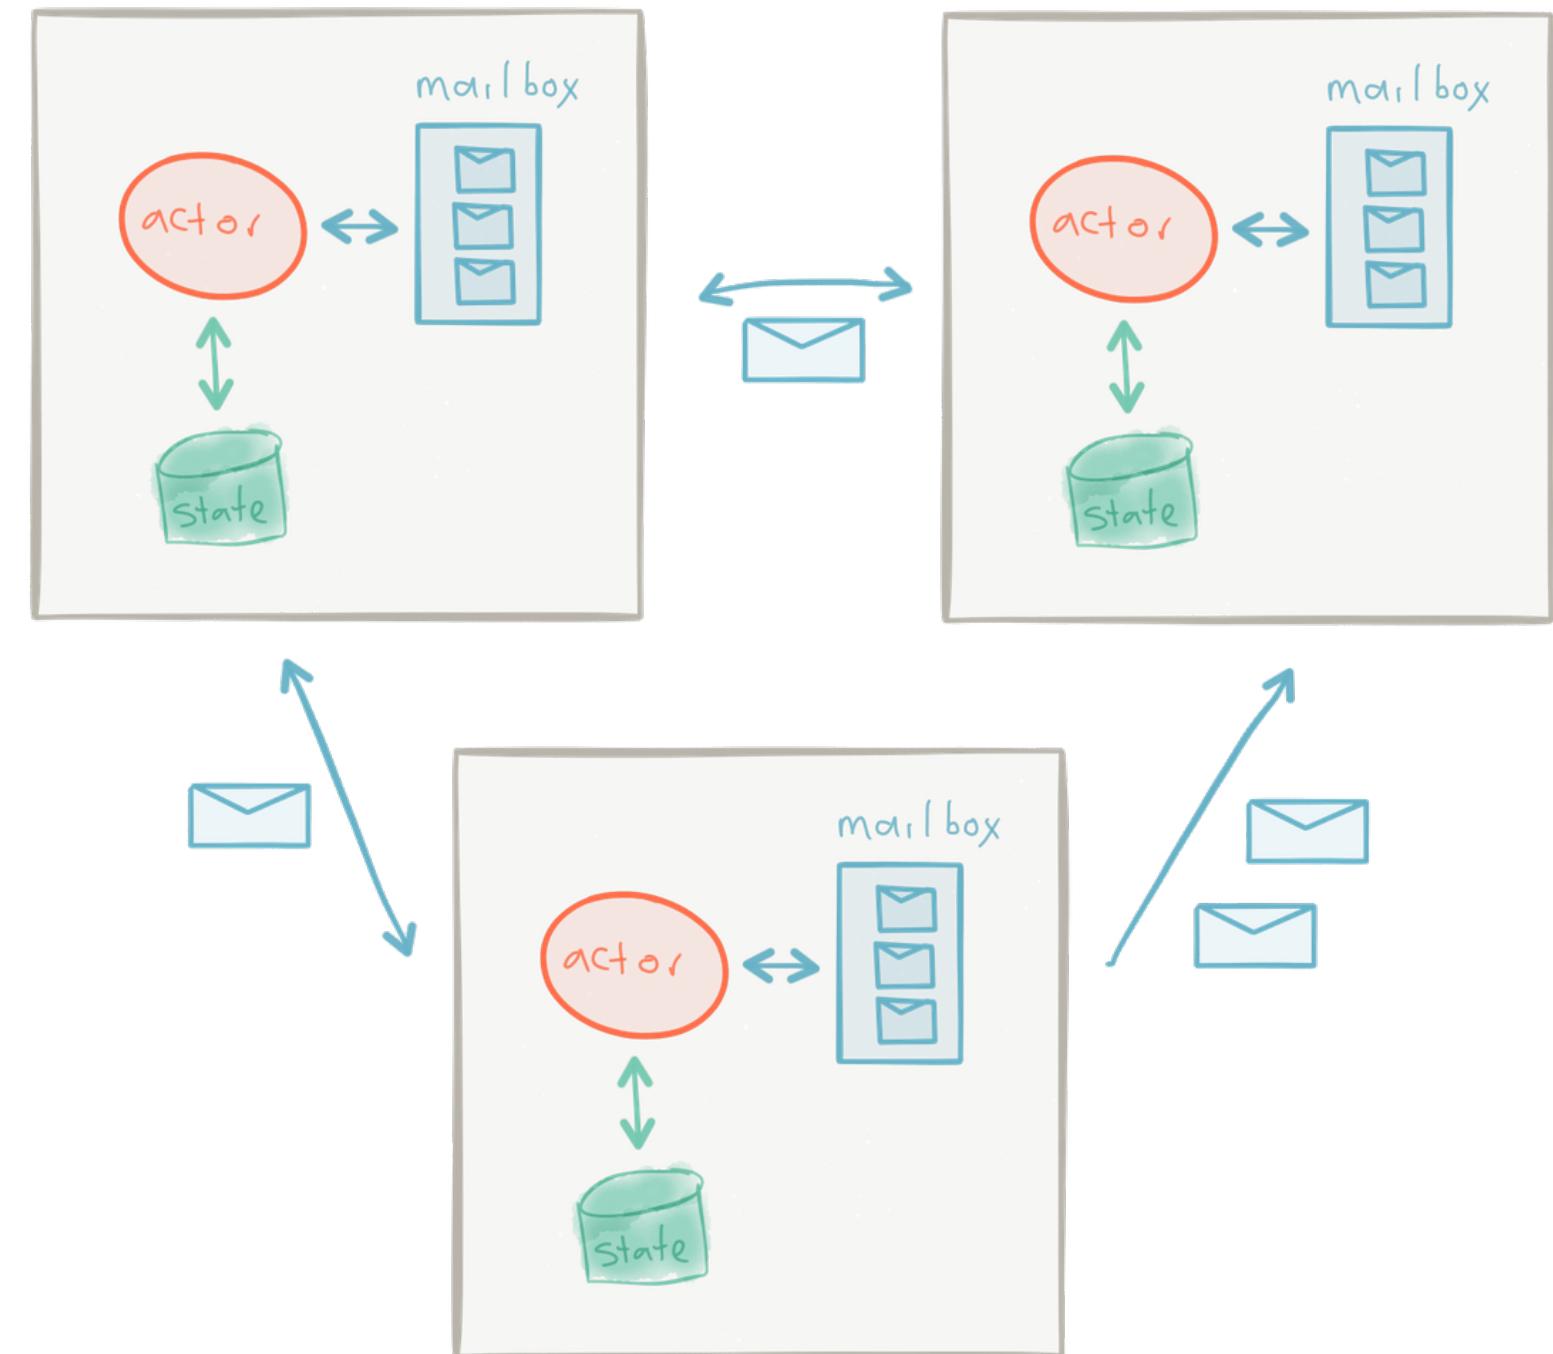
Considerations

- » Immutable, composable stream blueprints
- » Explicit materialization step
- » No magic at the expense of some extra code

MATERIALIZATION

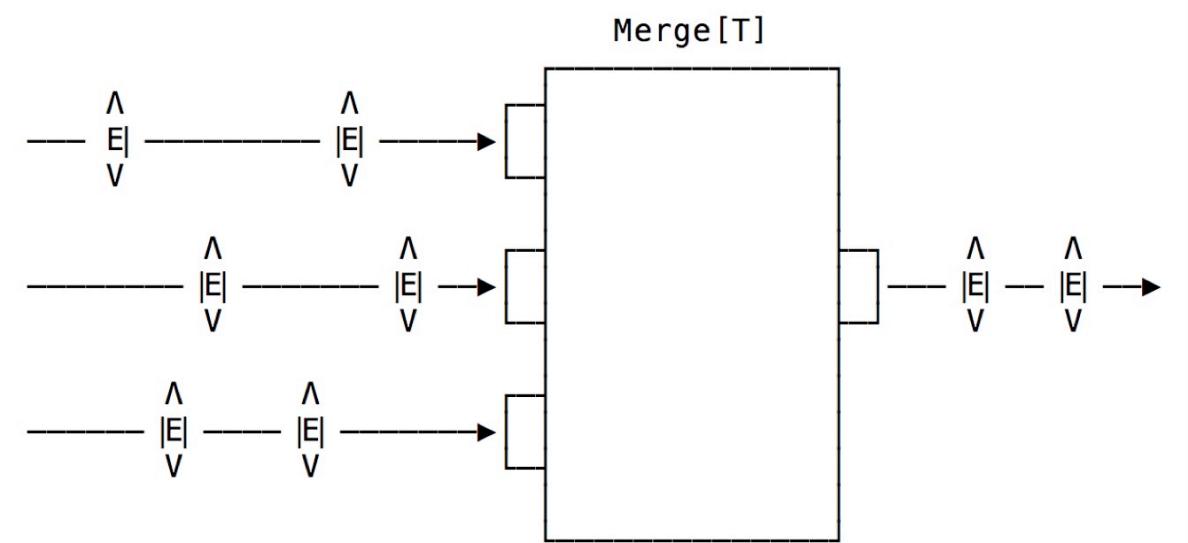
Separates the what from the how.

- » Use Source/Flow/Sink to create a blueprint
- » FlowMaterializer turns a blueprint into Akka Actors

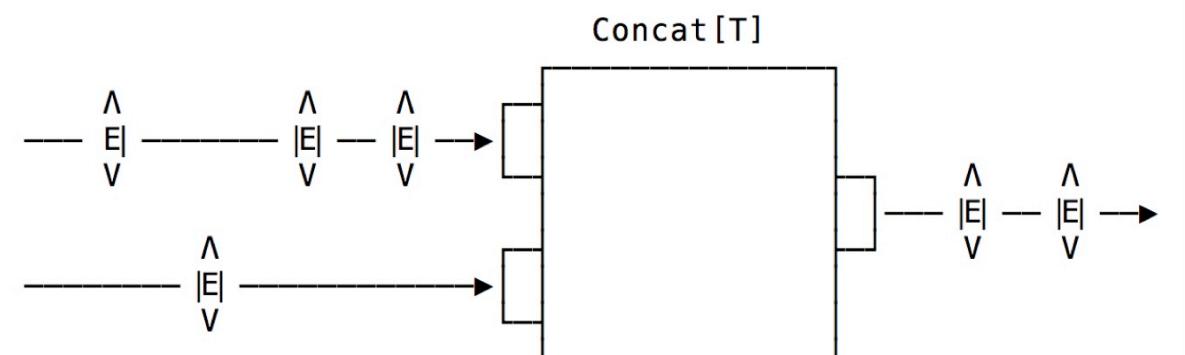
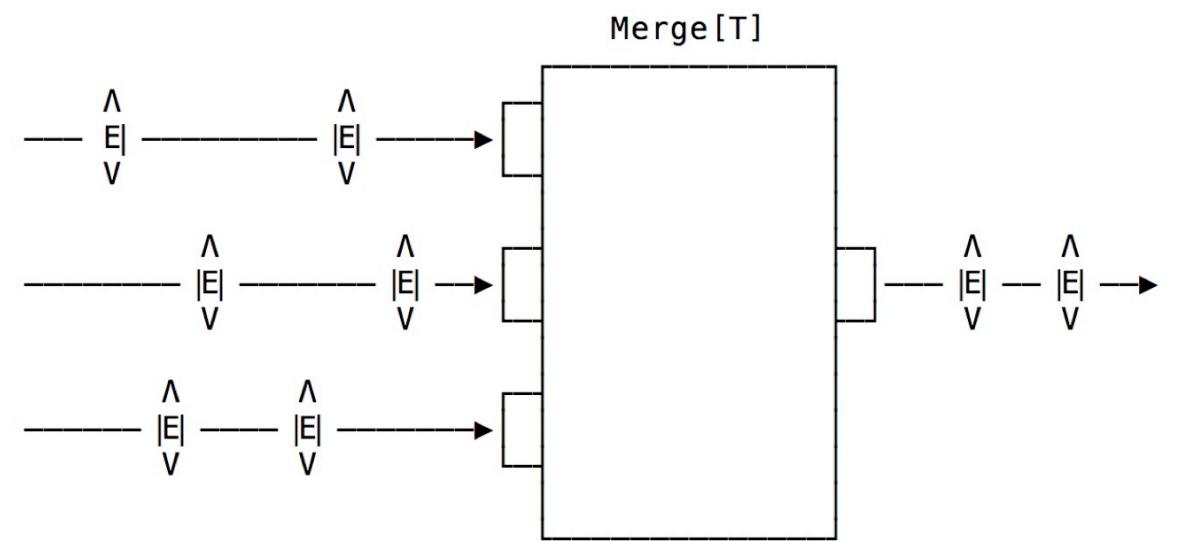


GRAPHS

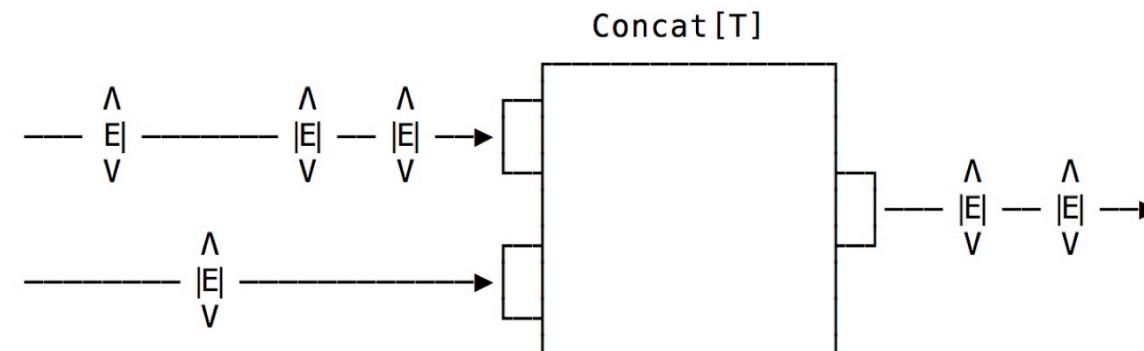
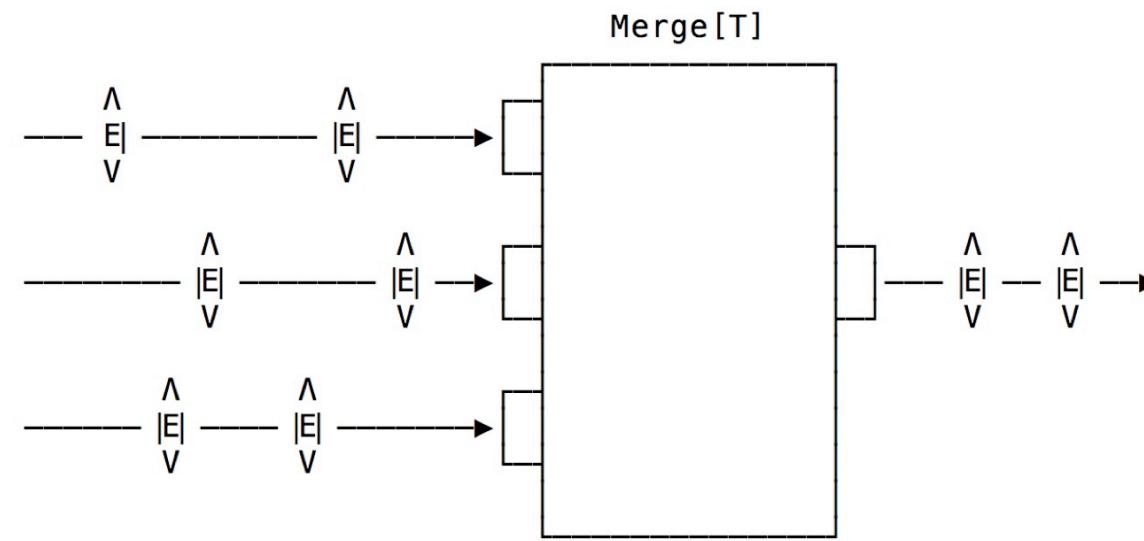
FAN IN - JOIN MULTIPLE STREAMS INTO A SINGLE STREAM



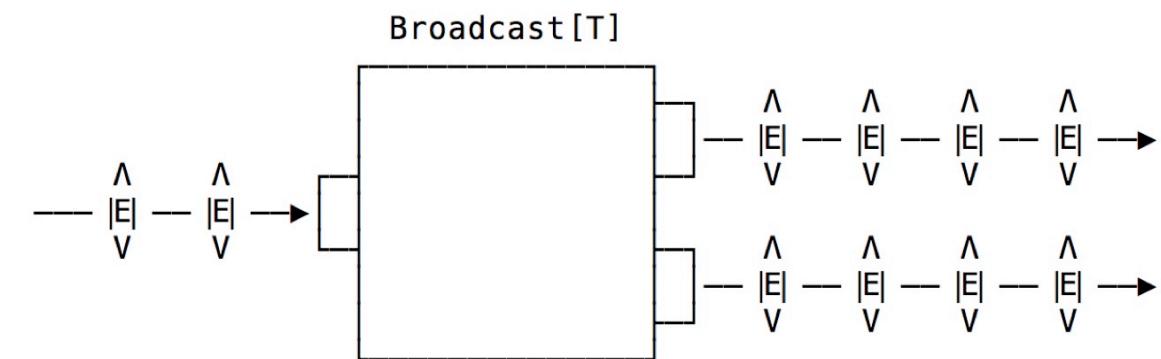
FAN IN - JOIN MULTIPLE STREAMS INTO A SINGLE STREAM



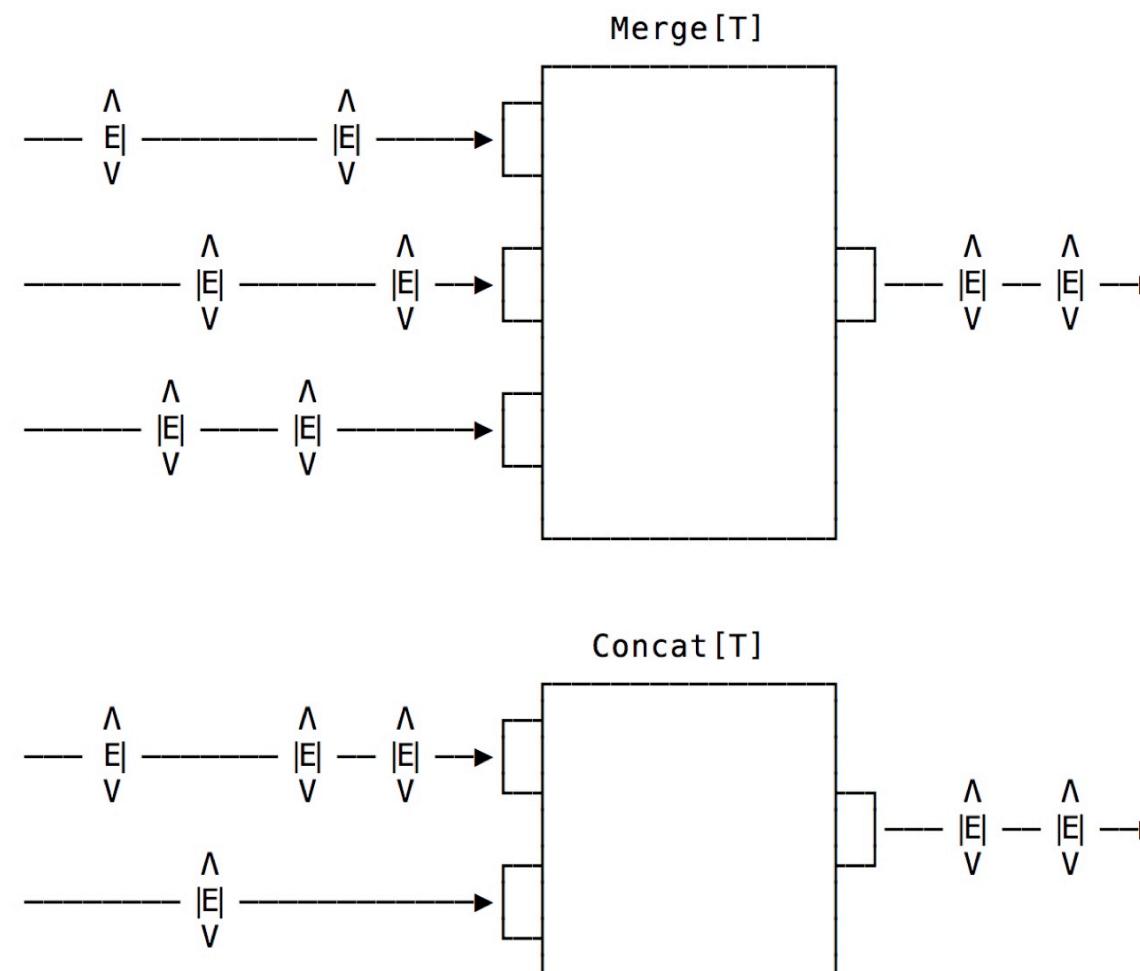
FAN IN - JOIN MULTIPLE STREAMS INTO A SINGLE STREAM



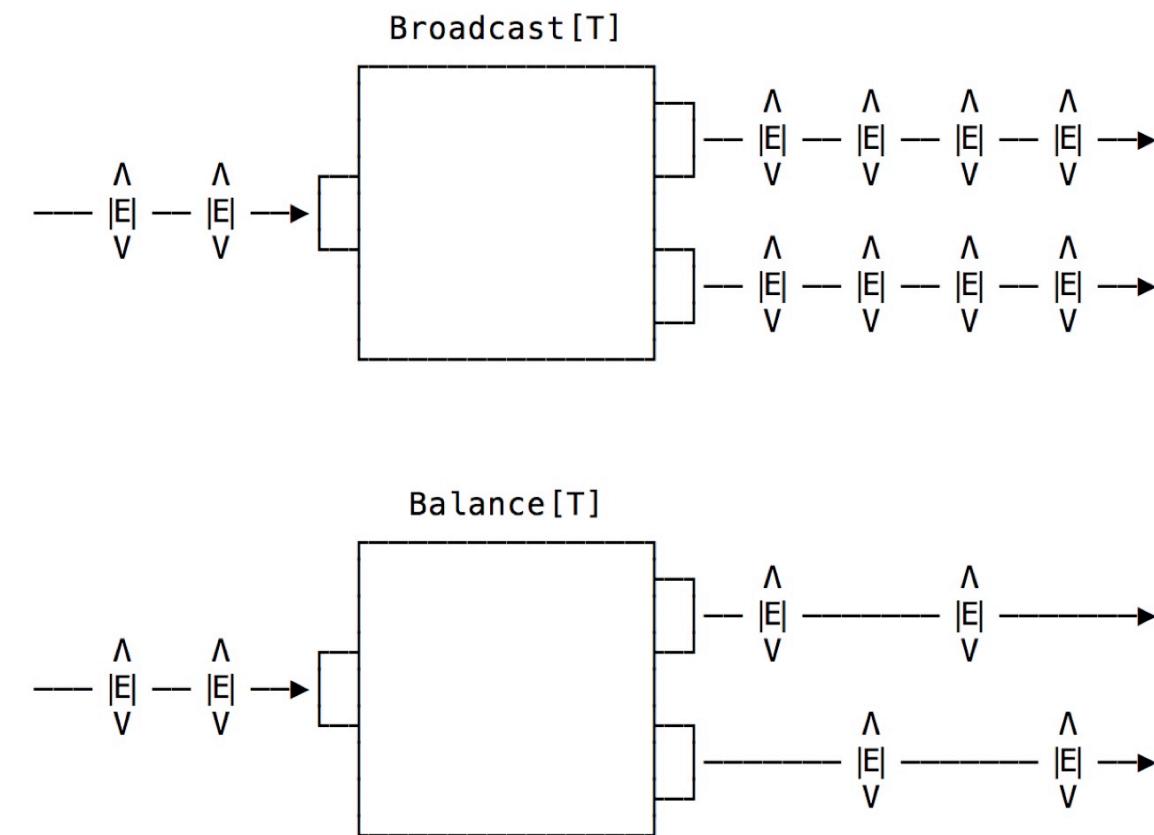
FAN OUT - SPLIT A STREAM INTO SUBSTREAMS



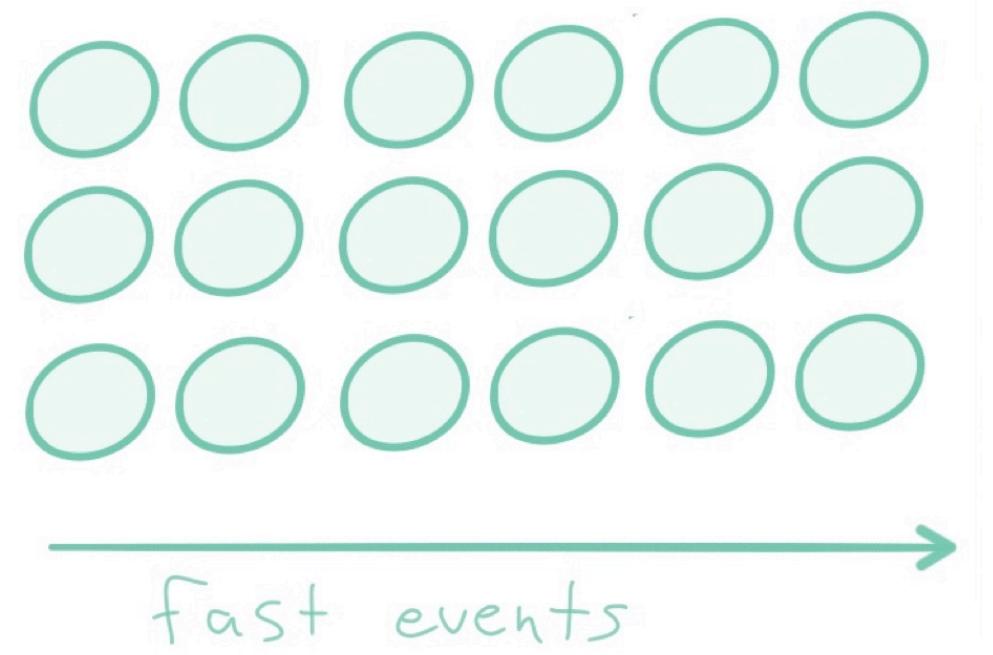
FAN IN - JOIN MULTIPLE STREAMS INTO A SINGLE STREAM



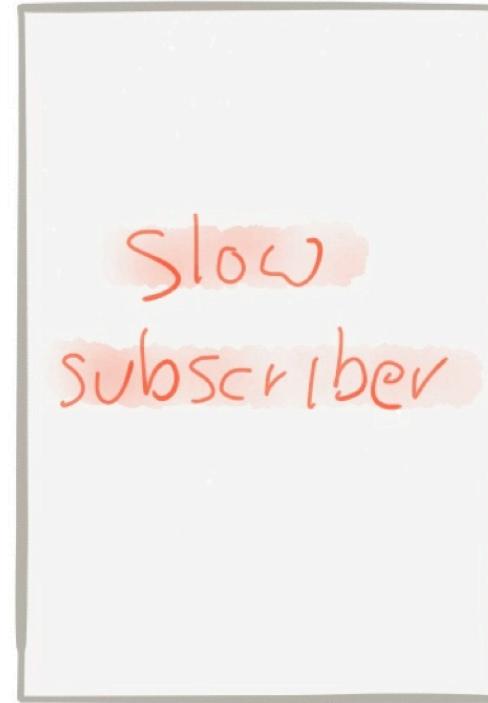
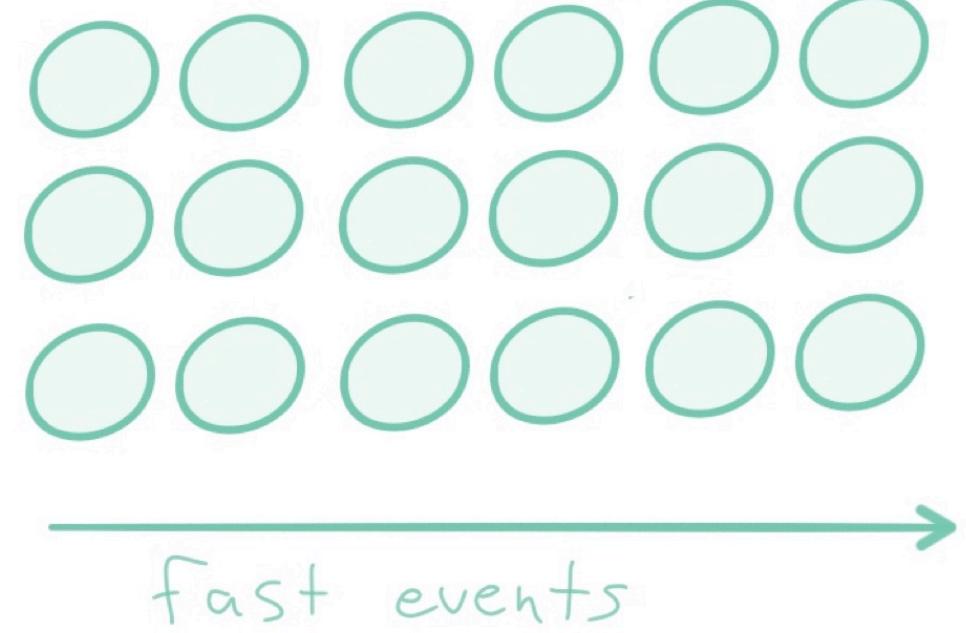
FAN OUT - SPLIT A STREAM INTO SUBSTREAMS

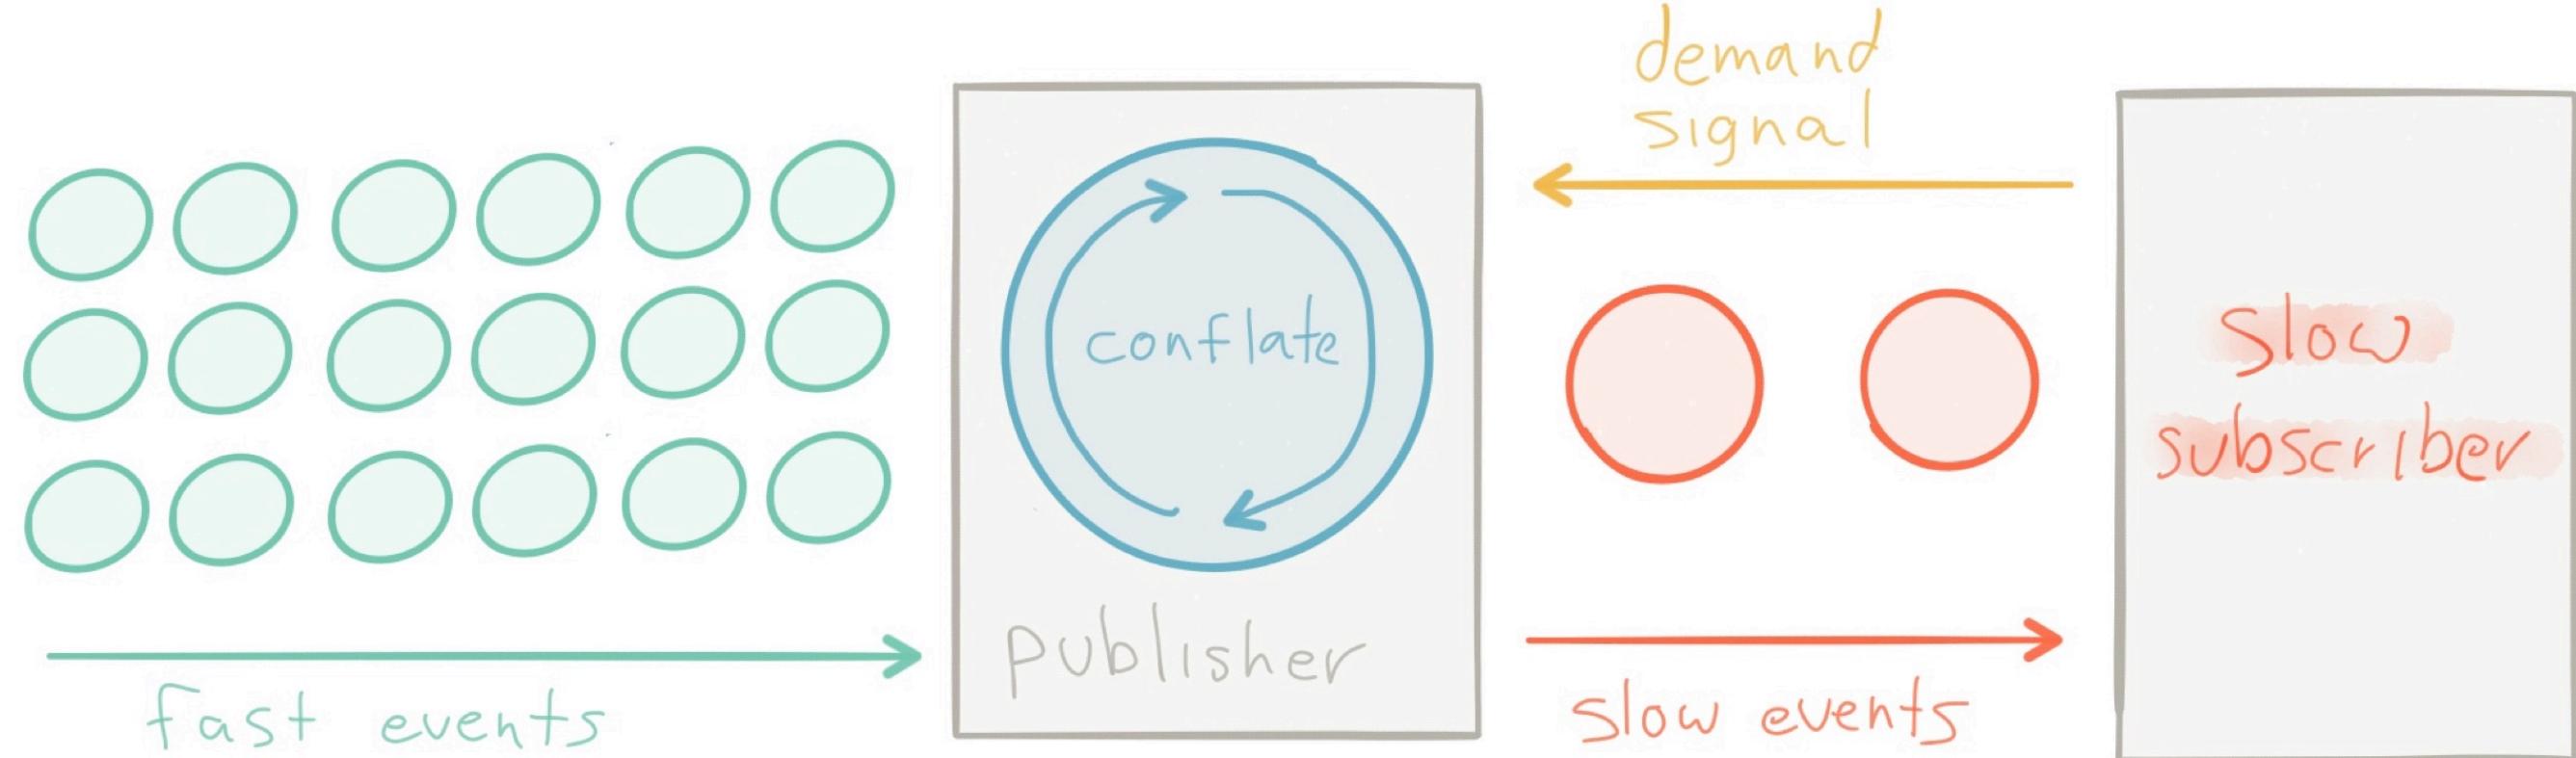


ADVANCED FLOW CONTROL

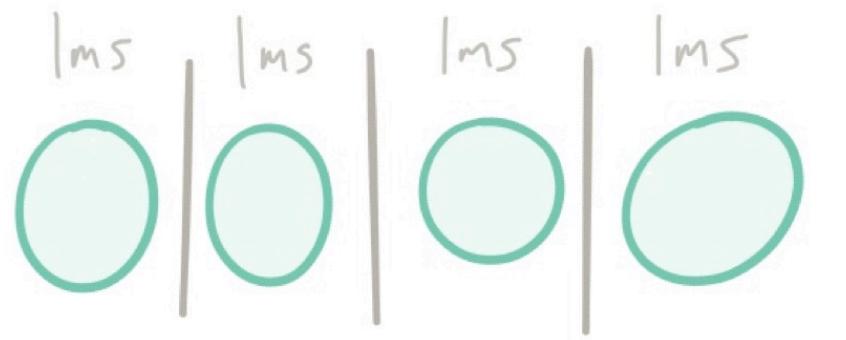


Slow
subscriber

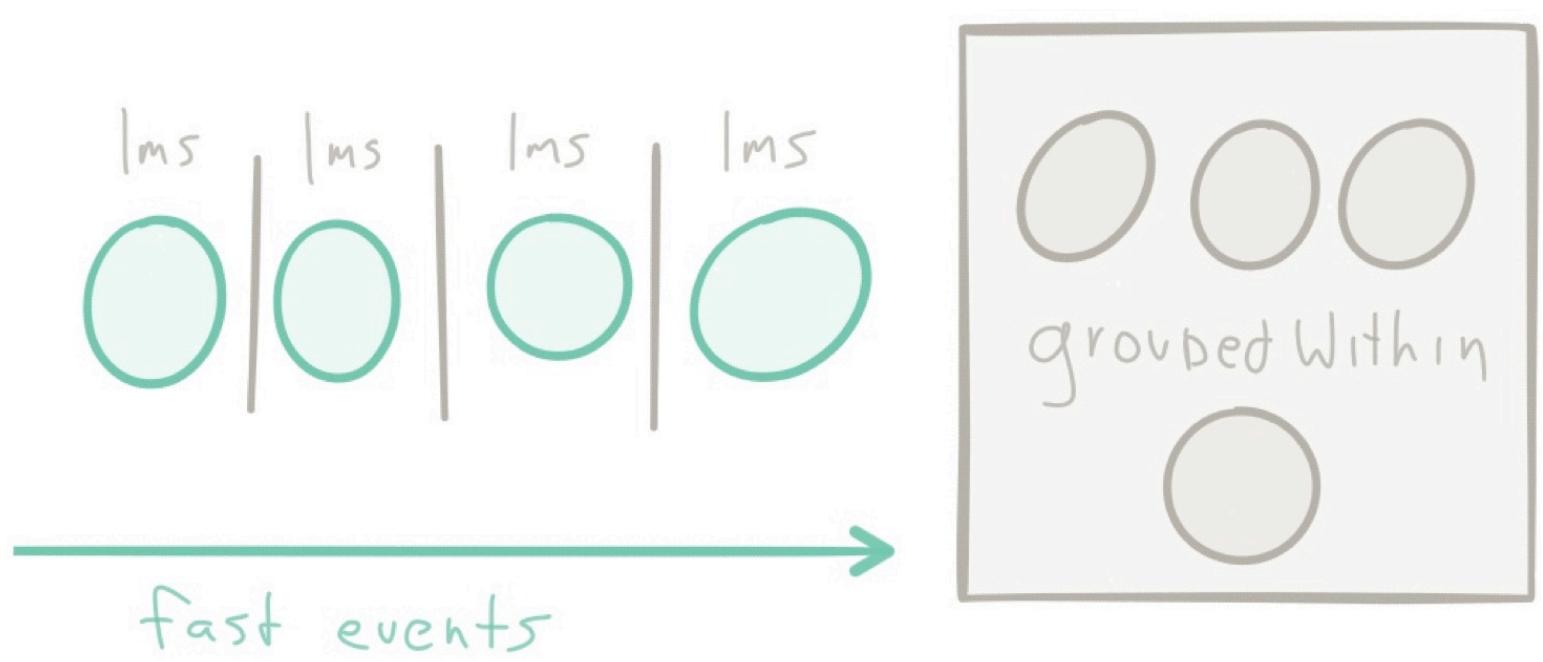


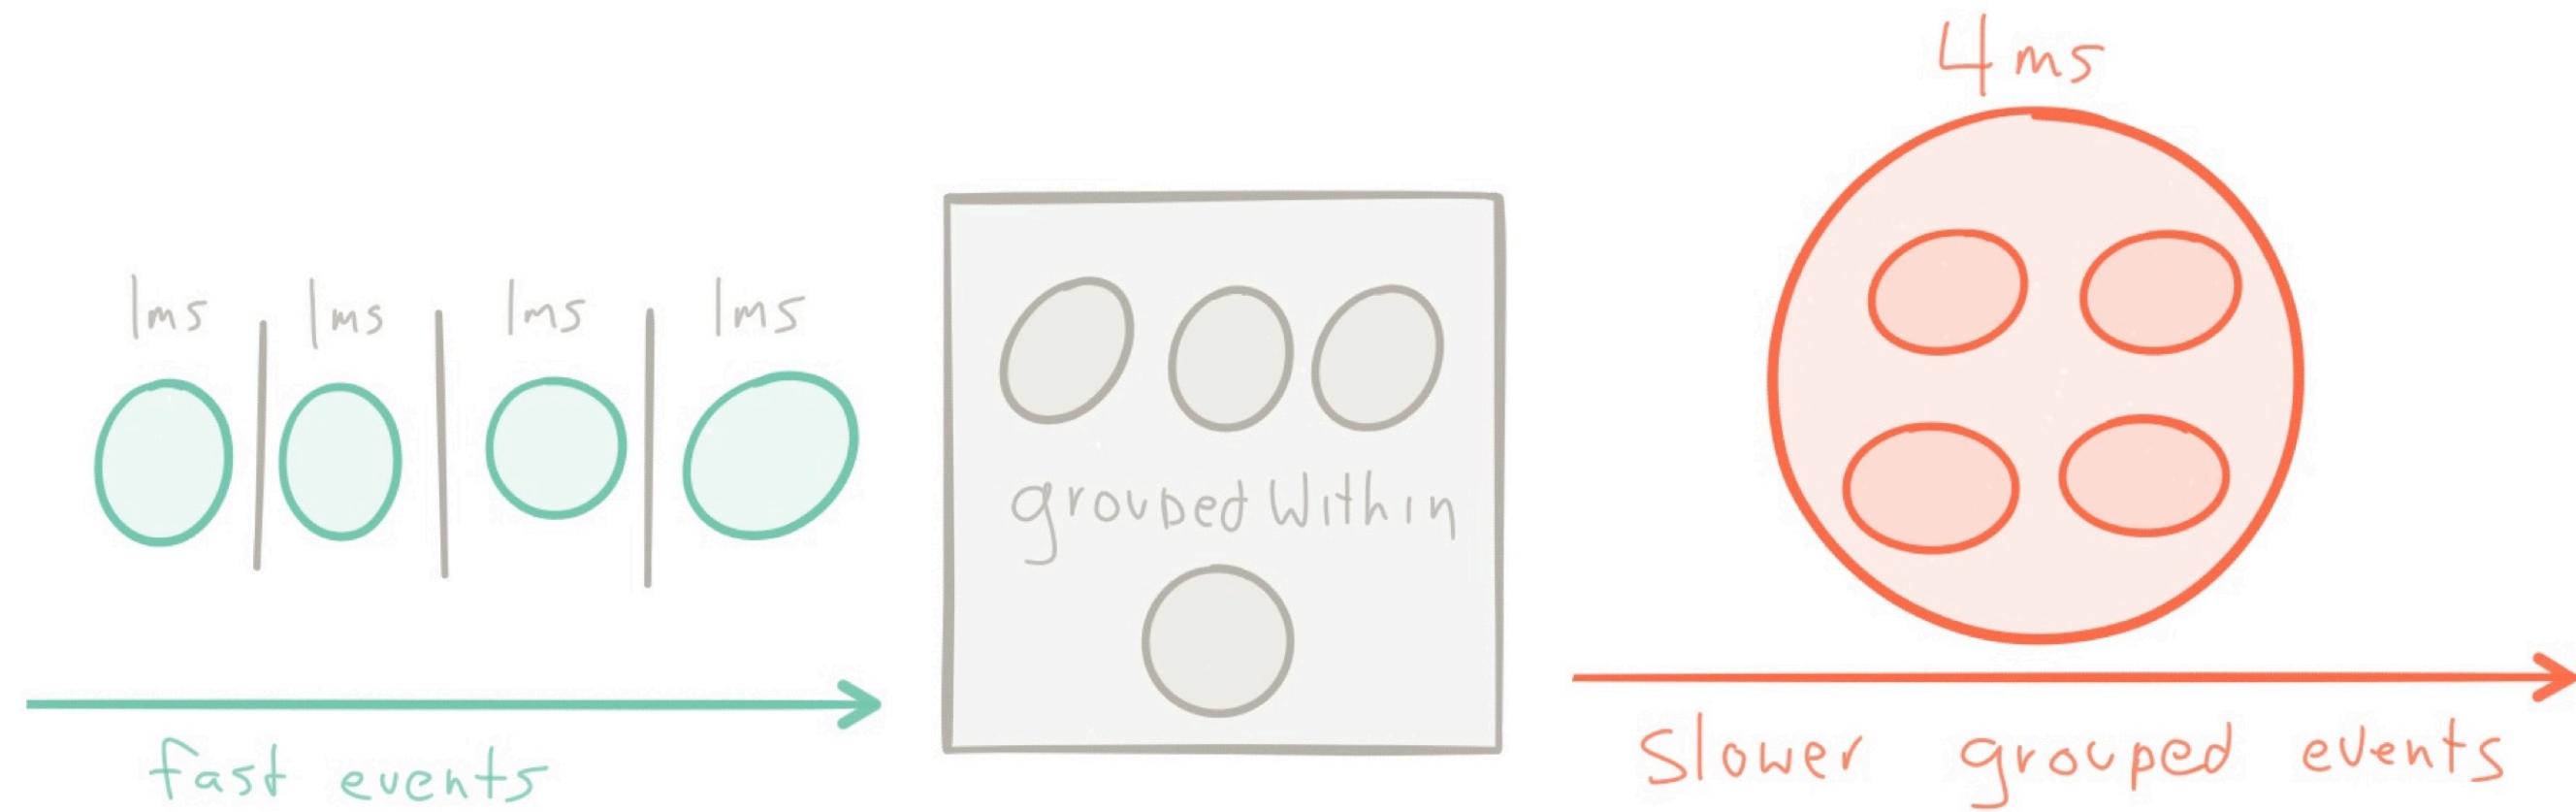


CONFLATE - EMIT AN AGGREGATE EVENT BASED ON ANY CRITERIA (AVERAGE, HIGH, LOW, ETC)



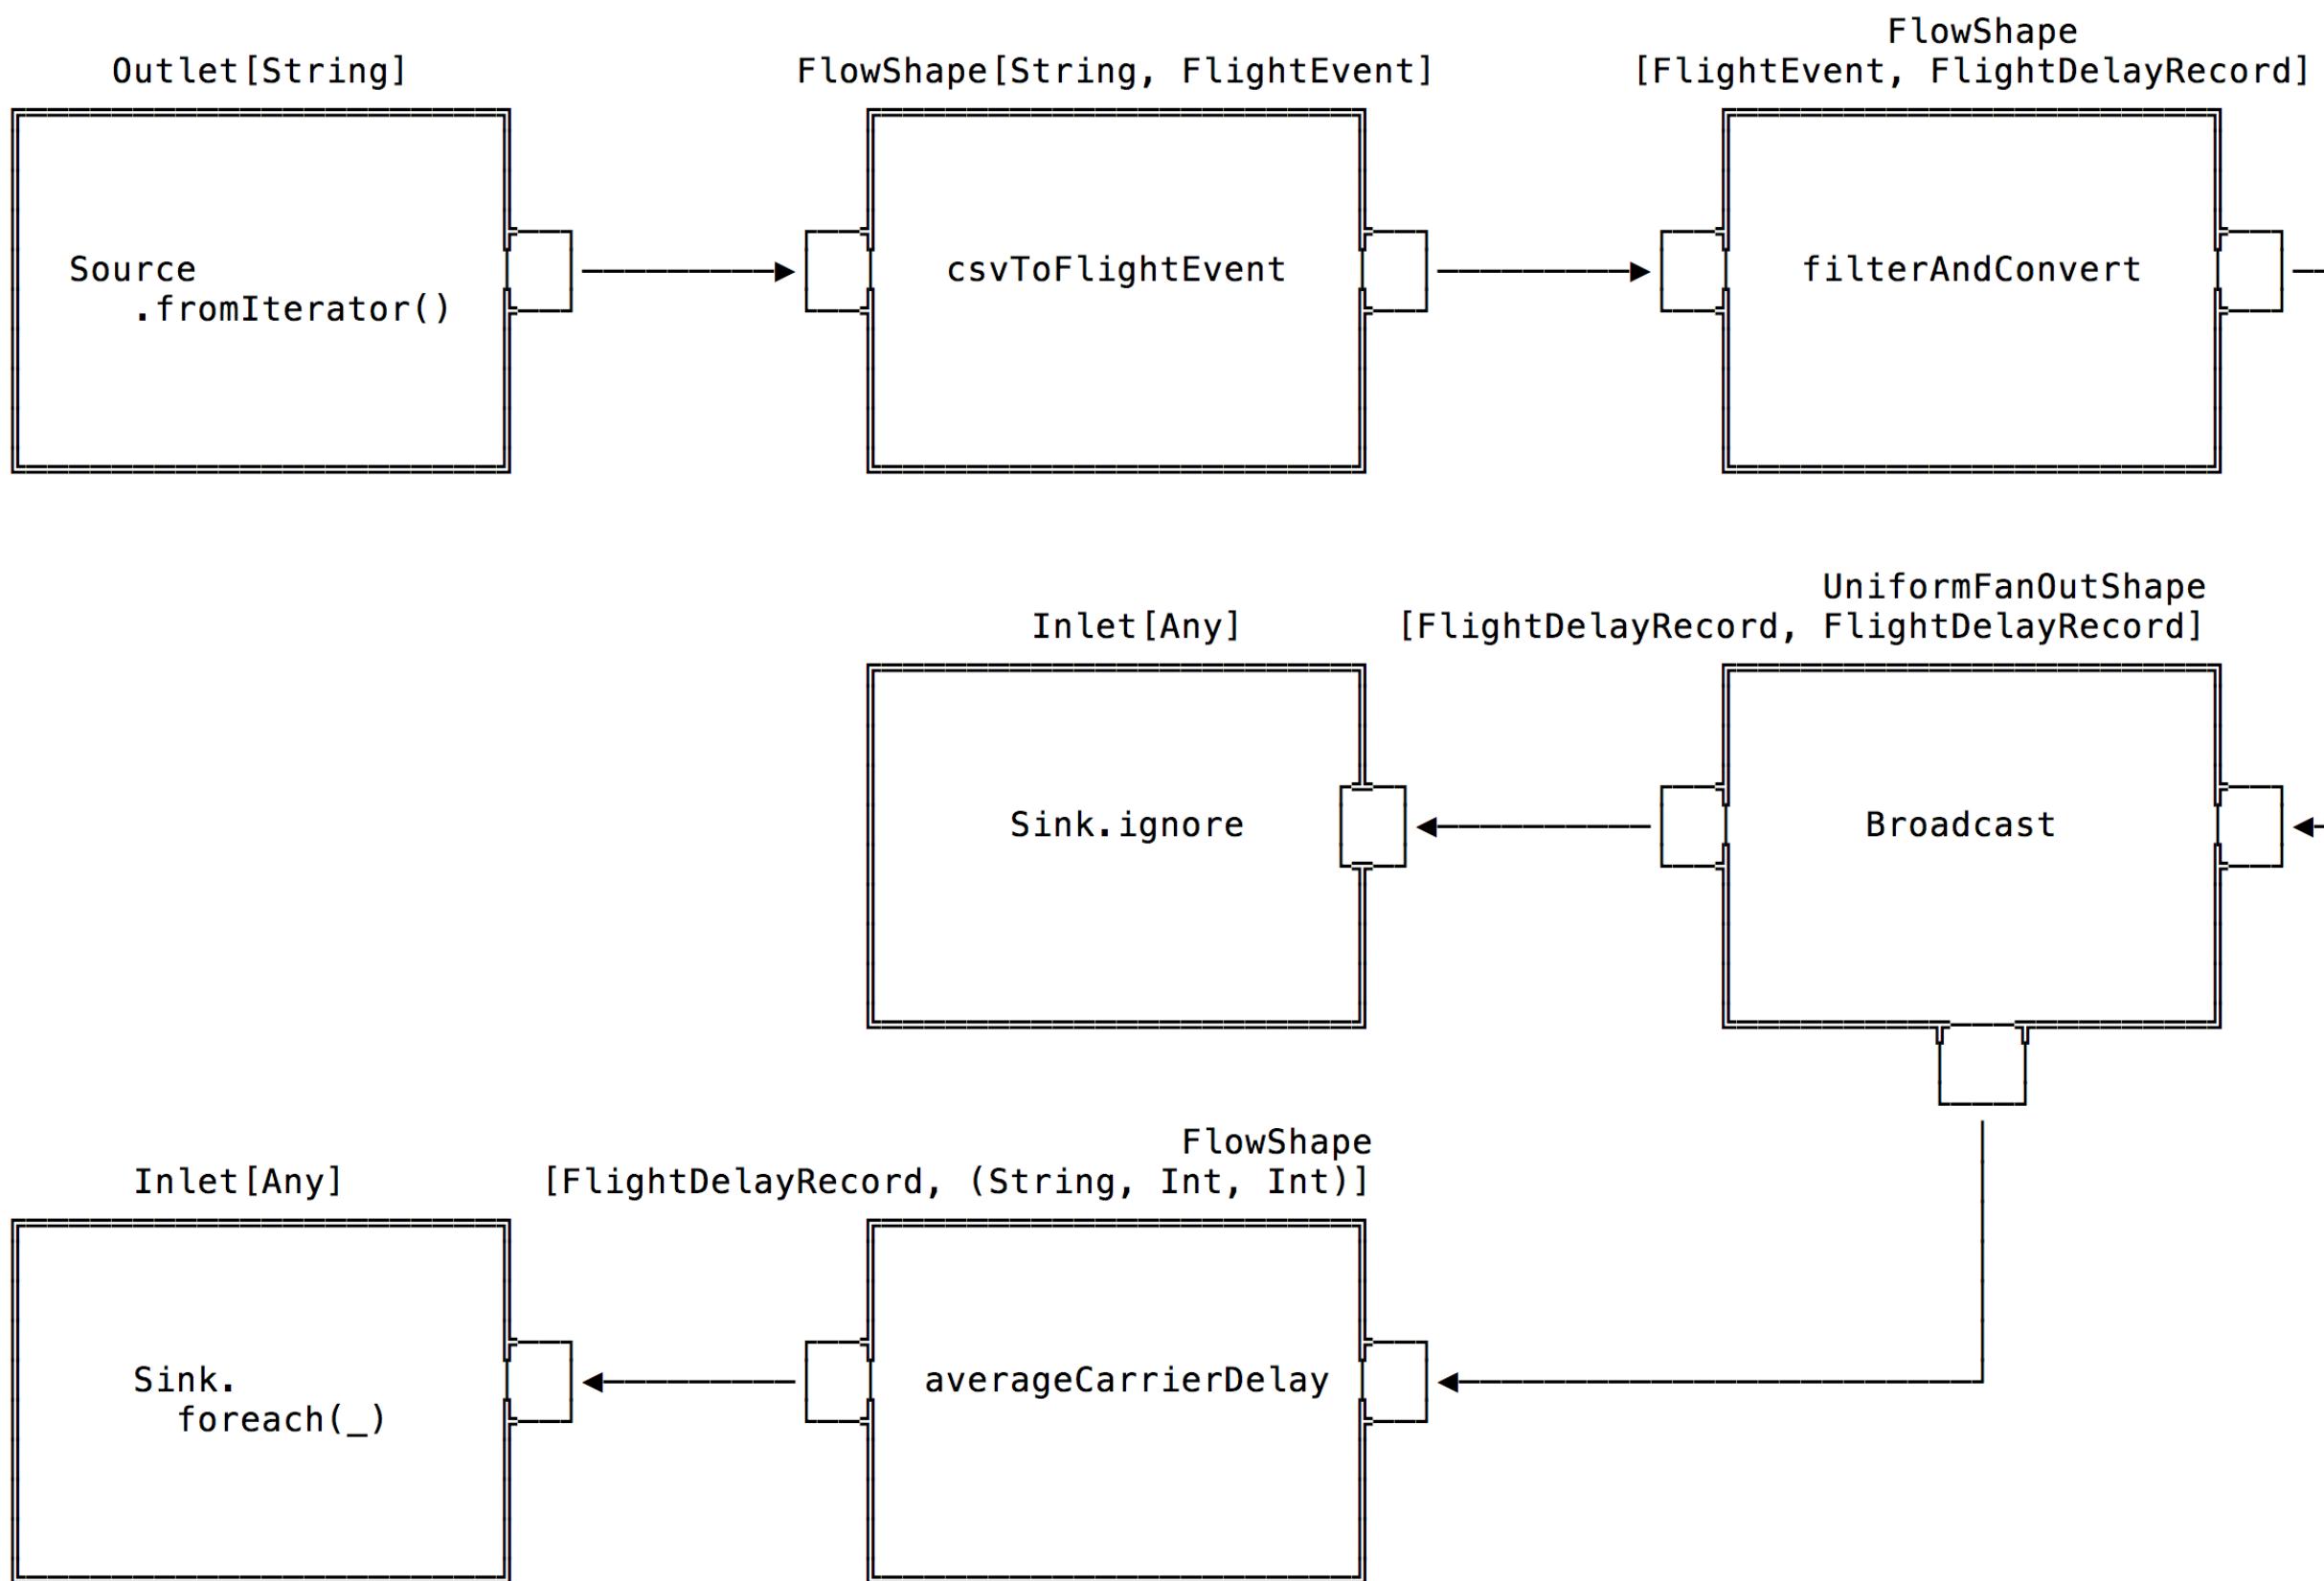
→
fast events





GROUPEDWITHIN - EMIT MINI-BATCH OF EVENTS BY TIME WINDOW OR # OF EVENTS

CODE AND REVIEW



```
val g = RunnableGraph.fromGraph(GraphDSL.create() {
    implicit builder =>
    import GraphDSL.Implicits._

    // Source
    val A: Outlet[String] = builder.add(Source.fromIterator(() => flightDelayLines)).out

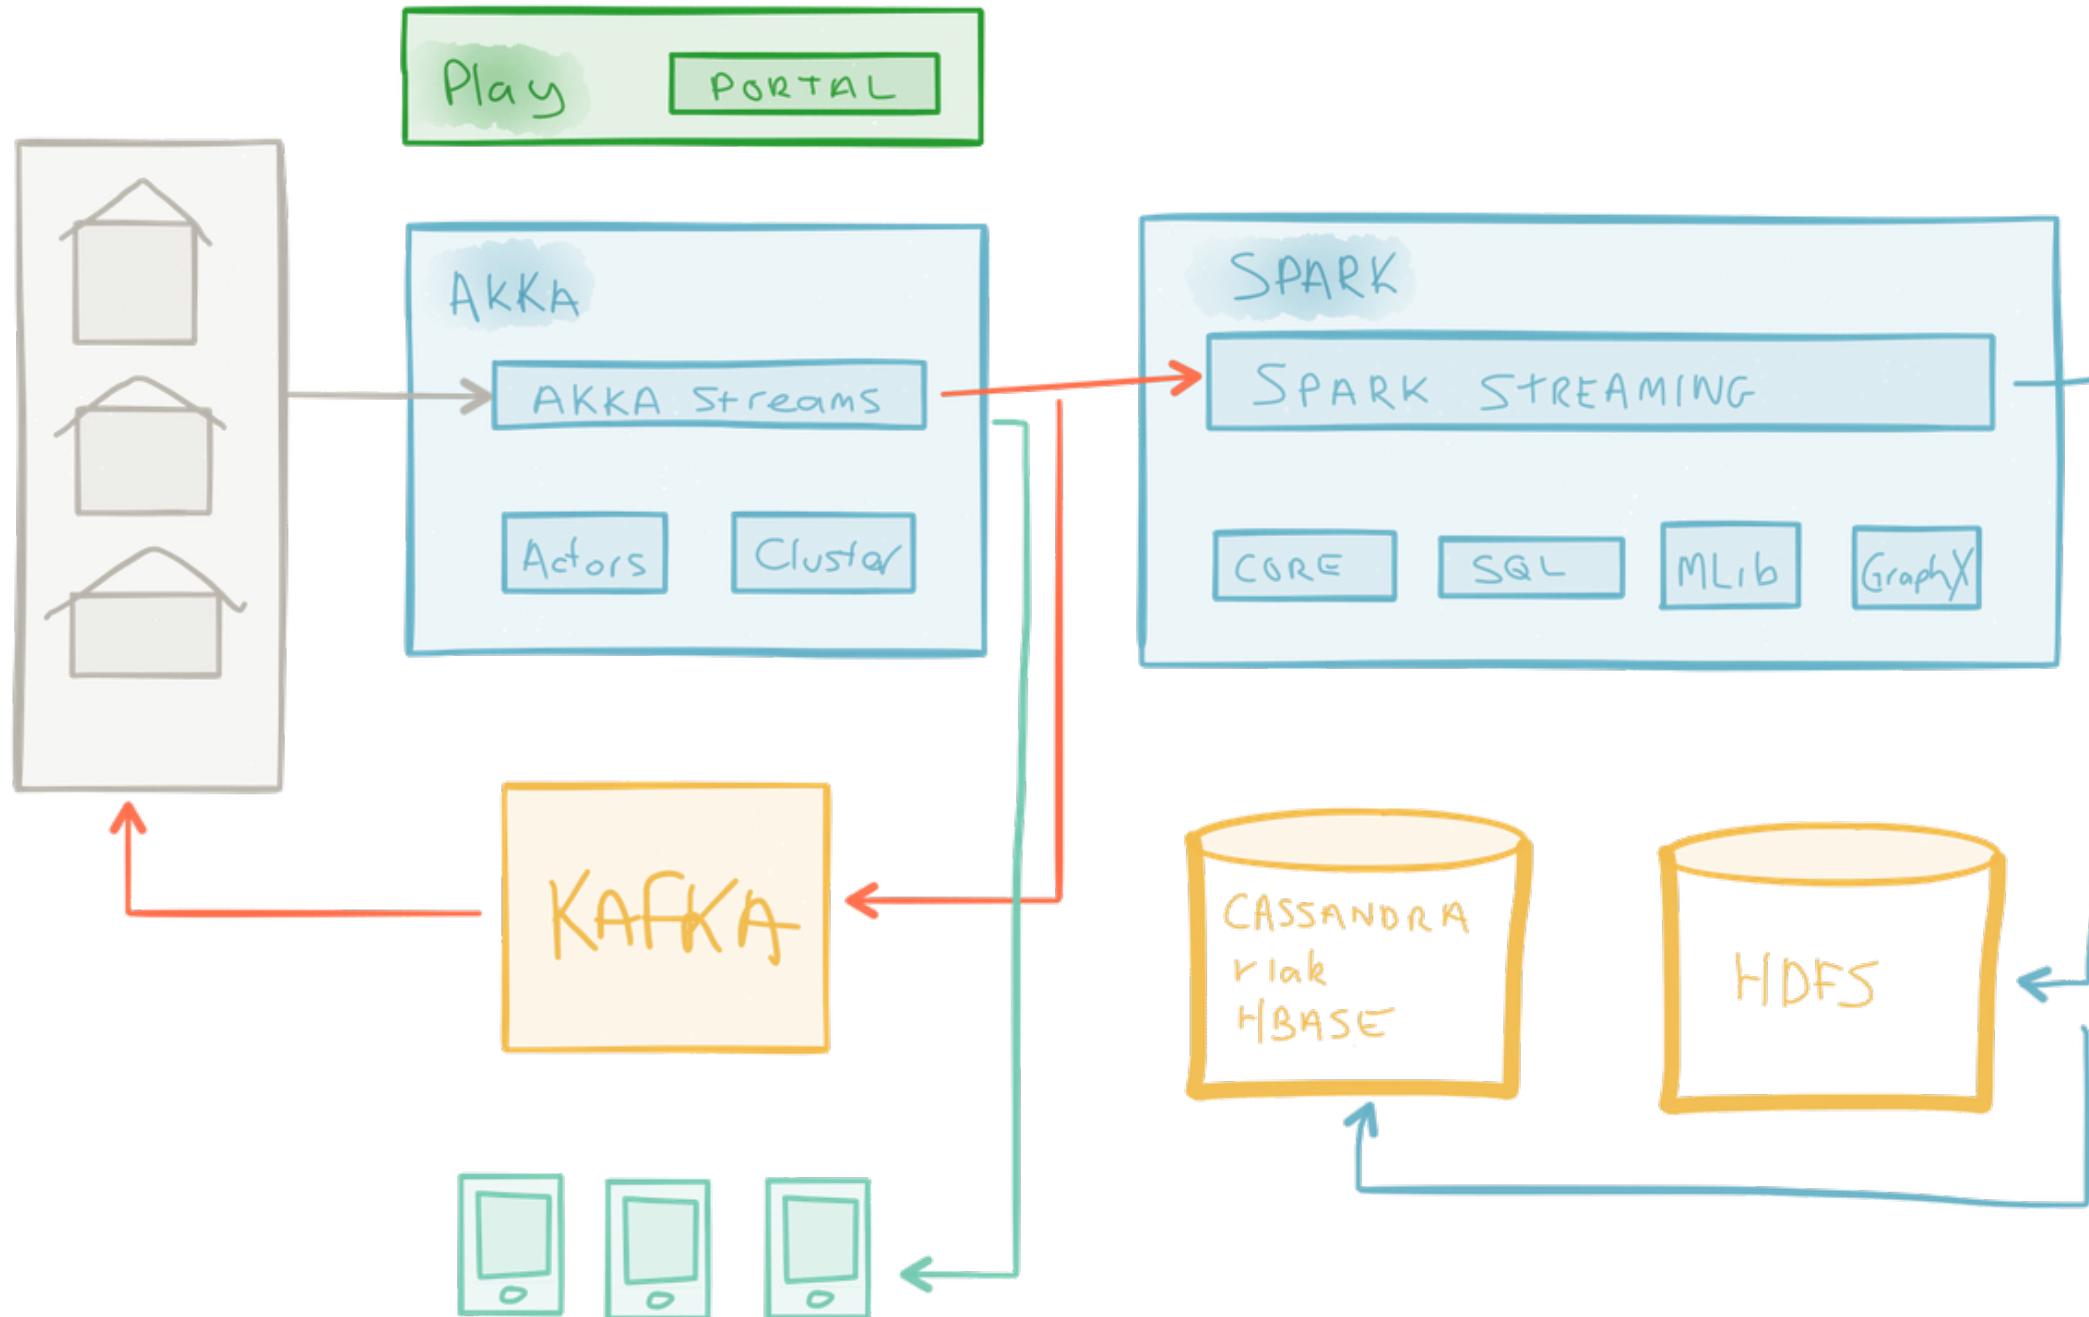
    // Flows
    val B: FlowShape[String, FlightEvent] = builder.add(csvToFlightEvent)
    val C: FlowShape[FlightEvent, FlightDelayRecord] = builder.add(filterAndConvert)
    val D: UniformFanOutShape[FlightDelayRecord, FlightDelayRecord] = builder.add(Broadcast[FlightDelayRecord](2))
    val F: FlowShape[FlightDelayRecord, (String, Int, Int)] = builder.add(averageCarrierDelay)

    // Sinks
    val E: Inlet[Any] = builder.add(Sink.ignore).in
    val G: Inlet[Any] = builder.add(Sink.foreach(averageSink)).in

    // Graph
    A ~> B ~> C ~> D
        E <~ D
    G <~ F <~ D

    ClosedShape
})

g.run()
```



THANK YOU!

Visit <https://www.lightbend.com/products/lightbend-reactive-platform> to get started

CONTACT INFO

- » Lightbend: <https://www.lightbend.com/contact>
- » Twitter: @kvnwbbbr
- » Email: kevin.webber@lightbend.com